

Convolutional Neural Networks

Spring 2018

Review

- Last week:
 - Neural Network: Gradient Calculation with Backpropagation
 - Neural Network: Weight Update Methods
 - Neural Network: Training
- Assignments (Canvas):
 - Lab assignment due Friday
 - Project pre-proposal due yesterday
 - Project proposal due next week
- Questions?

Lab 4: VizWiz Challenge Winners

- Class: # Correct Predictions from 30 Test Examples
 - 15
 - 17 (x2)
 - 18 (x4)
 - 19 (x2)
 - 20 (x3)
 - 21 (x2)
 - 26 (x1)
- First place
 - Linli Ding (26/30)
- Second Place
 - Jennie Wolfgang (21/20)
 - Xiaomeng Mu (21/20)

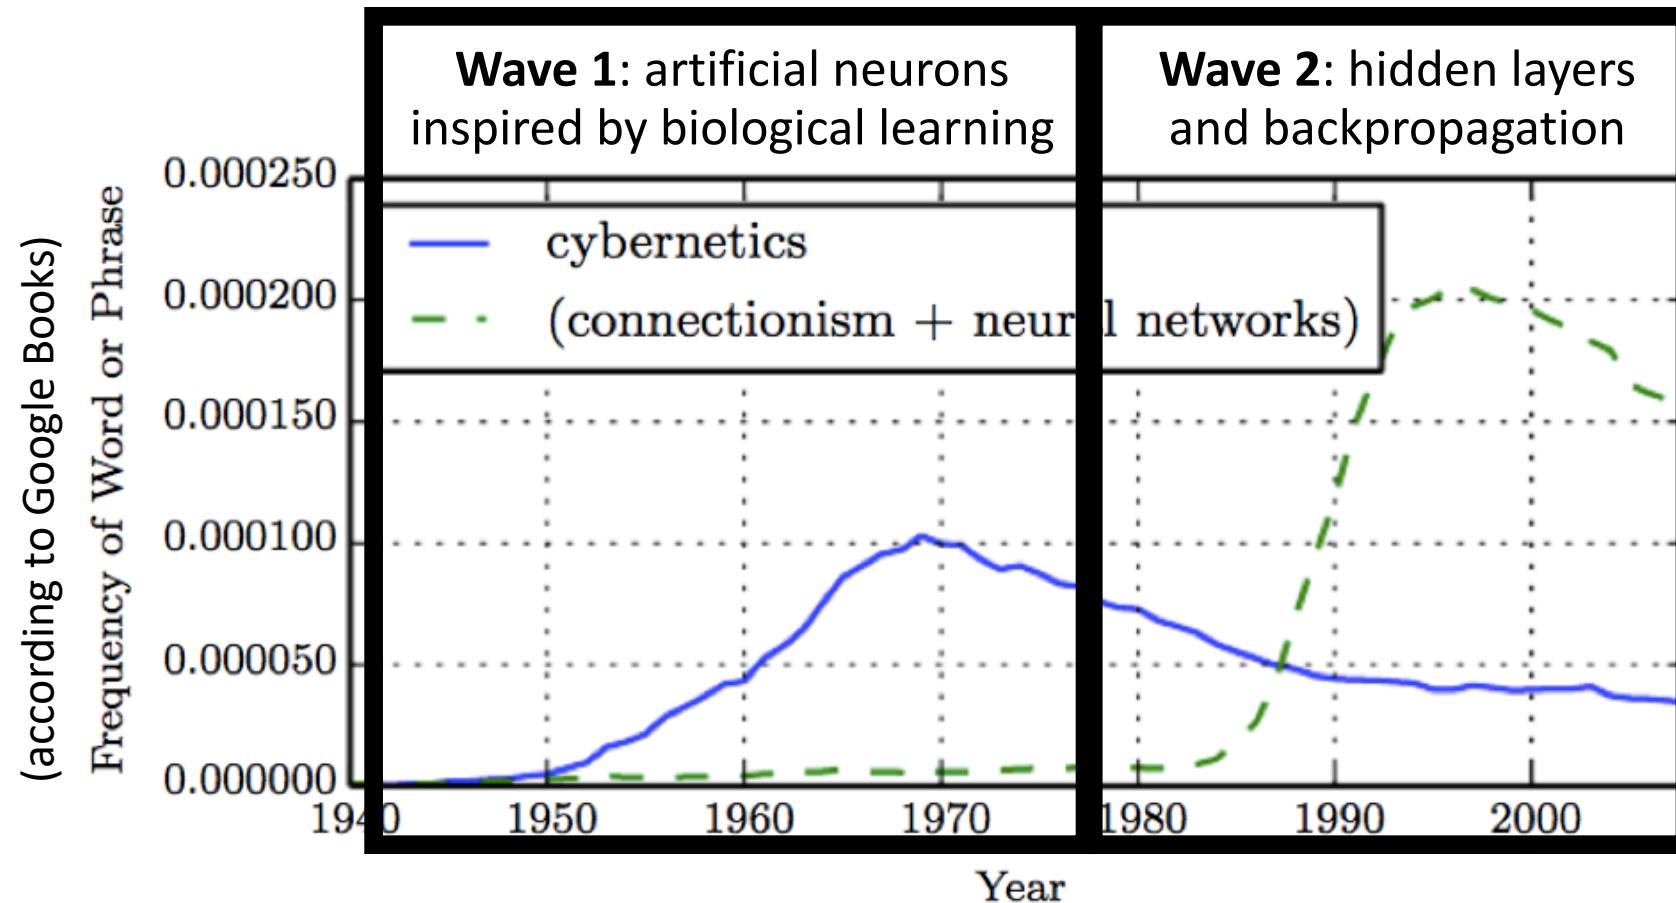
Today's Topics

- Revisiting History of Modern Neural Networks and “Deep Learning”
- Convolutional Neural Networks (CNN)
- Deep Features
- Lab

Today's Topics

- Revisiting History of Modern Neural Networks and “Deep Learning”
- Convolutional Neural Networks (CNN)
- Deep Features
- Lab

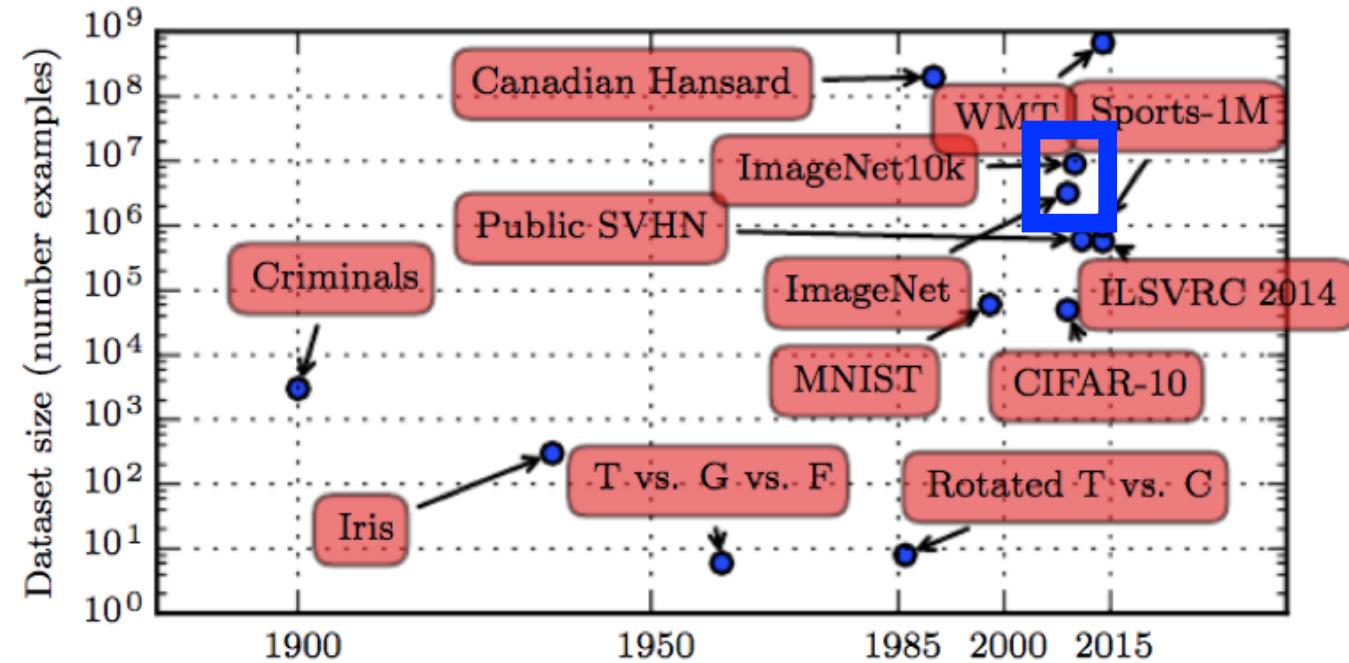
The History of “Deep Learning”



Wave 3: field emerged again as “deep learning” in 2006

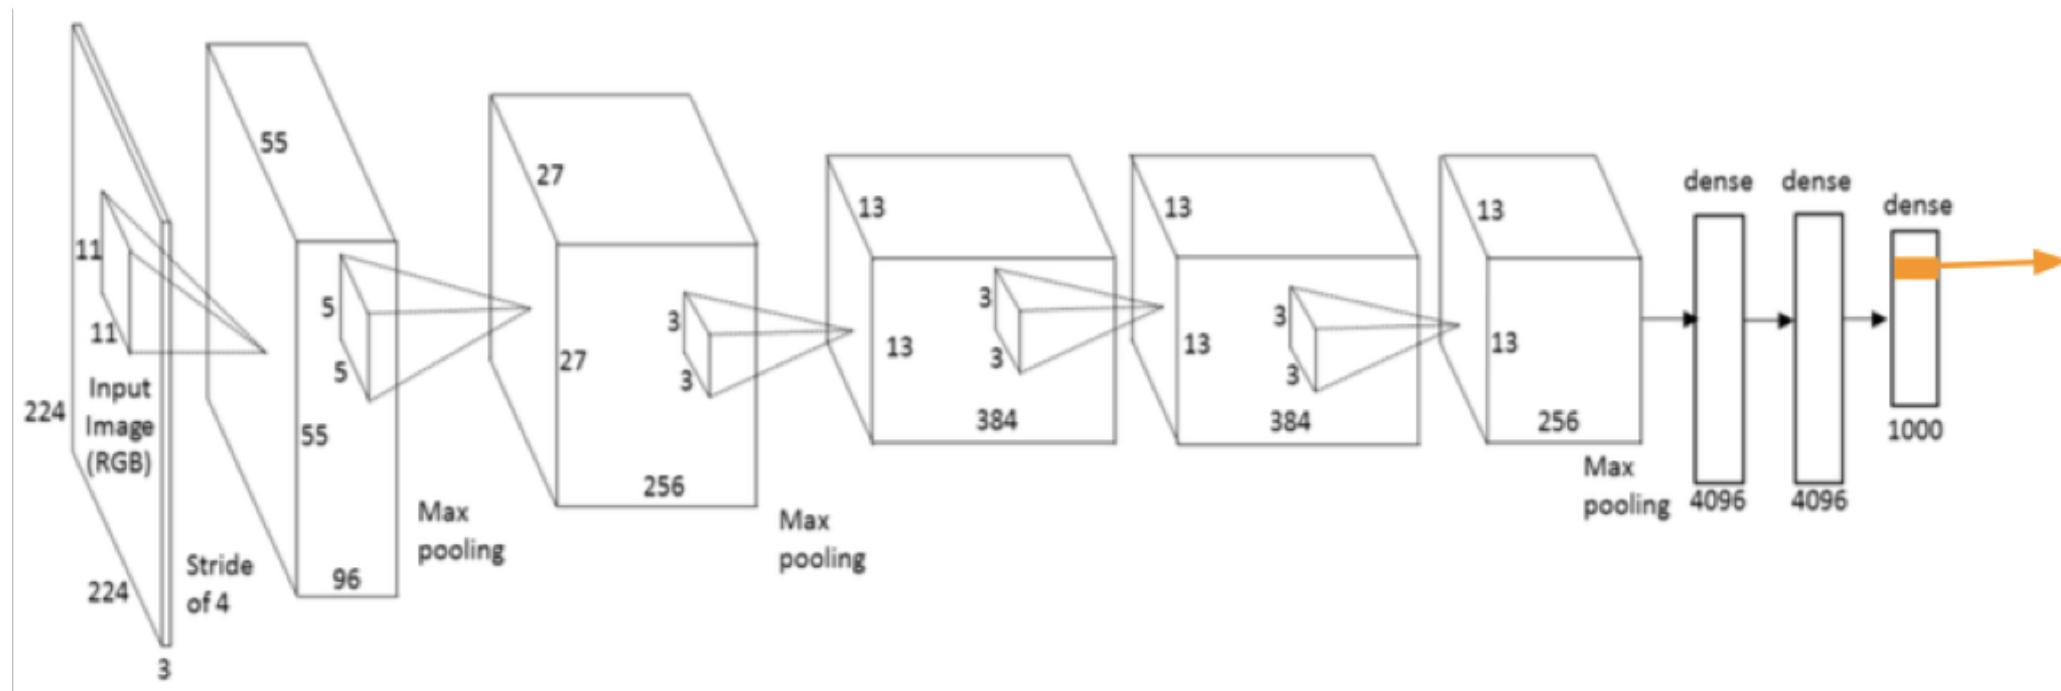
Deep Learning: What Prompted the 3rd Wave?

- The ability to provide the algorithms with the resources they need:
 1. “Big data”; e.g.,
 2. Faster Hardware; e.g., GPUs
 3. More memory; e.g., RAM



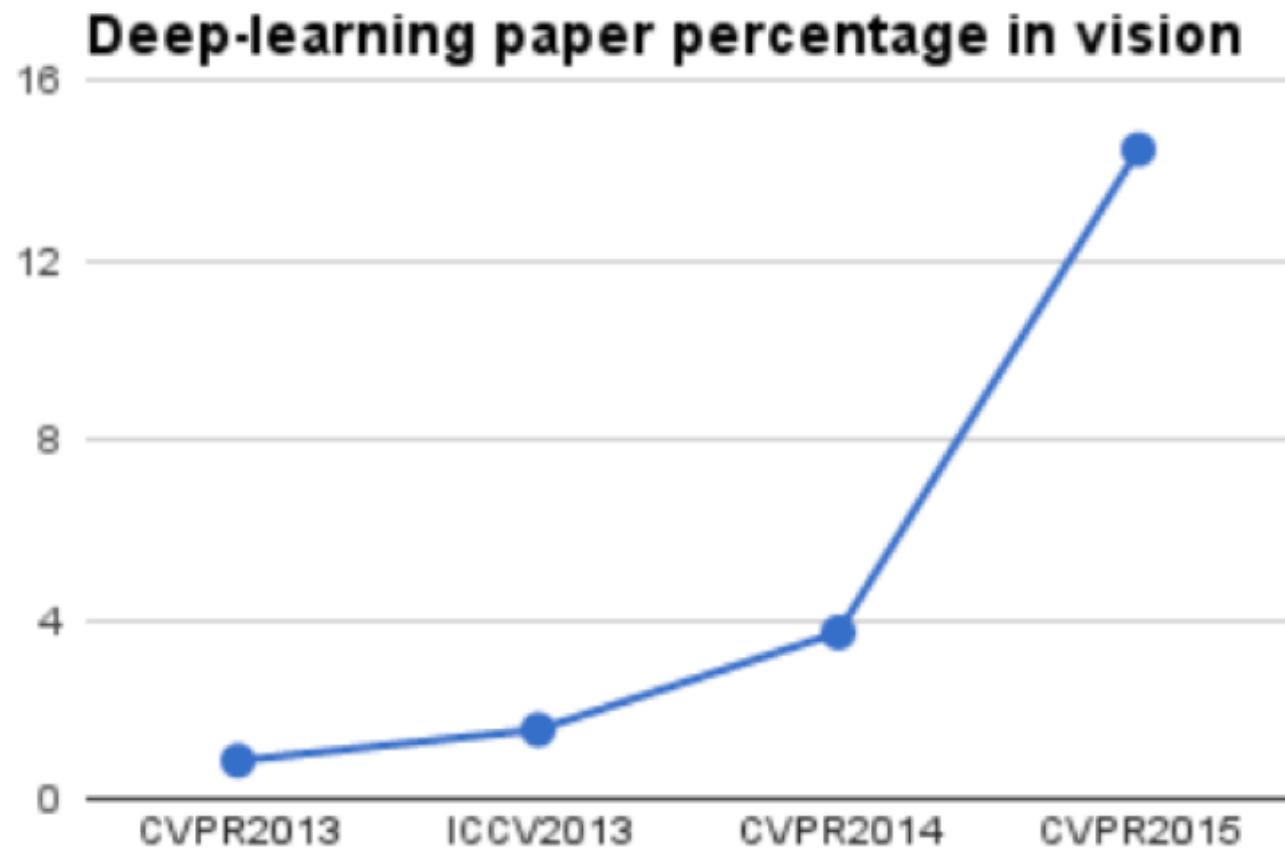
Why is Deep Learning So Popular Today?

- 2012: a **Convolutional Neural Network** won **ImageNet challenge**
 - state-of-the-art top-5 error rate fell from **26.1%** to **15.3%** with “AlexNet”



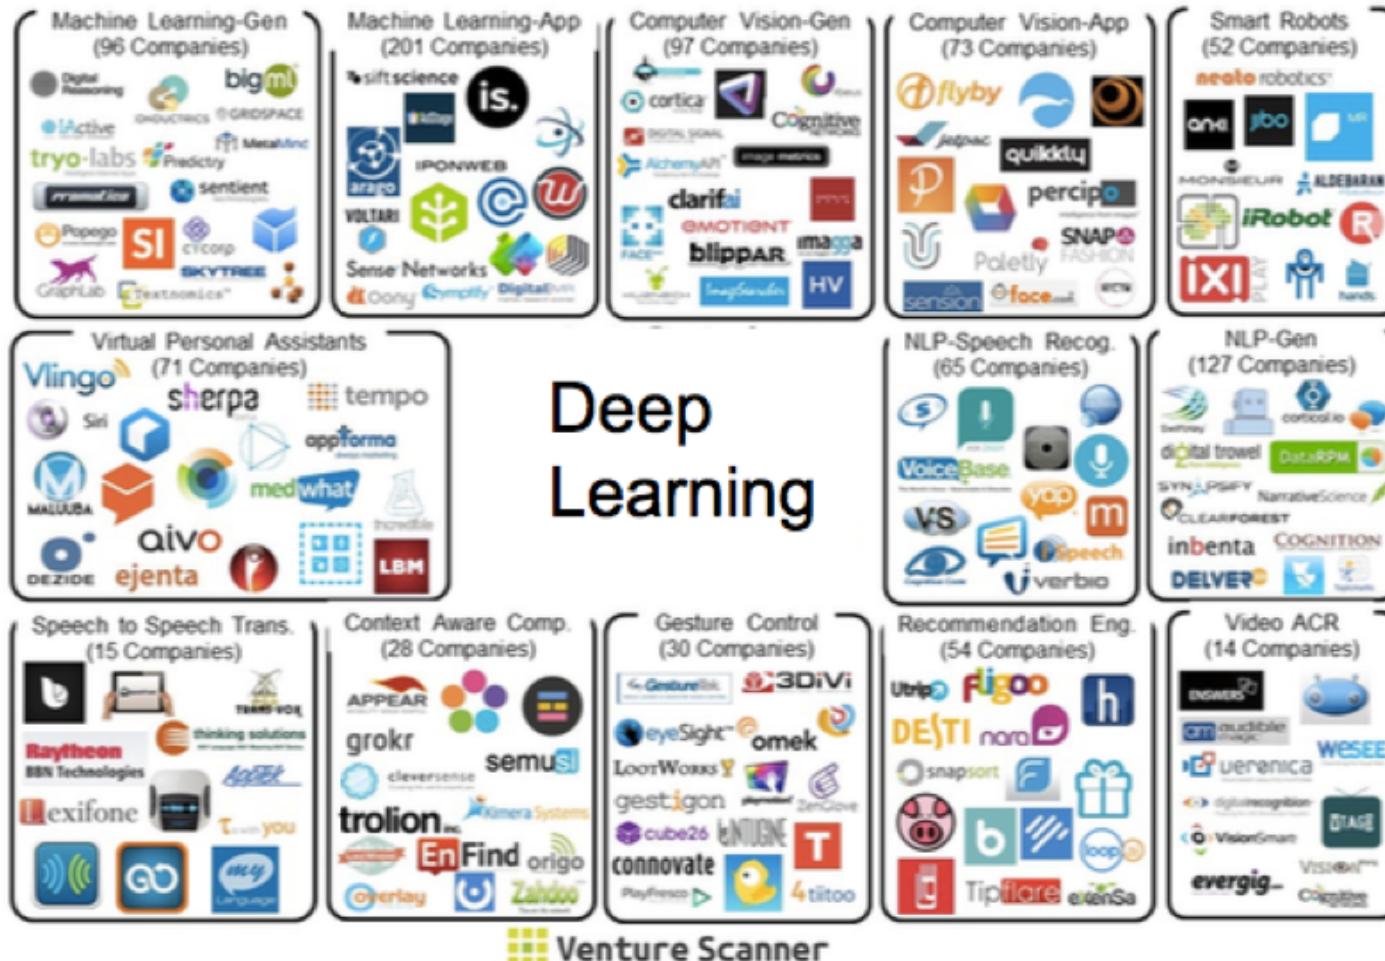
Slide Credit: <https://www.slideshare.net/xavigiro/saliency-prediction-using-deep-learning-techniques>
A. Krizhevsky, I. Sutskever, G. E. Hinton “ImageNet classification with deep convolutional neural networks”

Why is Deep Learning So Popular Today?



Slide Credit: <https://www.slideshare.net/xavigiro/saliency-prediction-using-deep-learning-techniques>

Deep Learning Inspiring New Companies



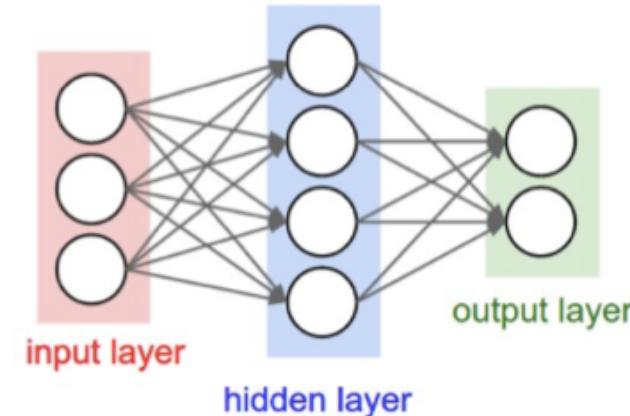
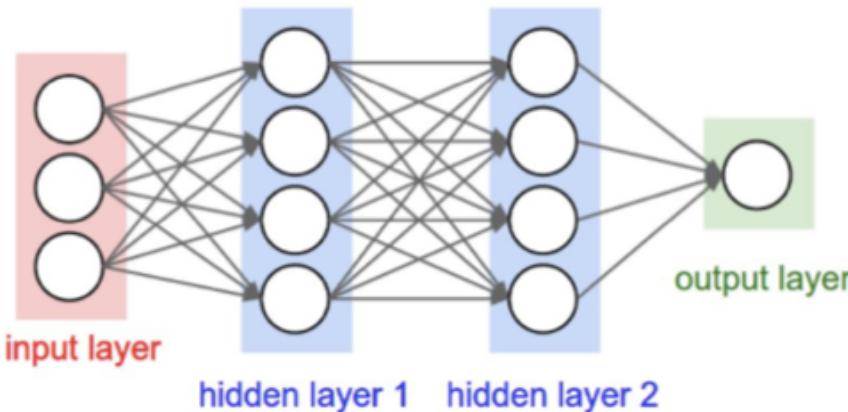
Slide Credit: <https://www.slideshare.net/xavigiro/saliency-prediction-using-deep-learning-techniques>

Today's Topics

- Revisiting History of Modern Neural Networks and “Deep Learning”
- Convolutional Neural Networks (CNN)
- Deep Features
- Lab

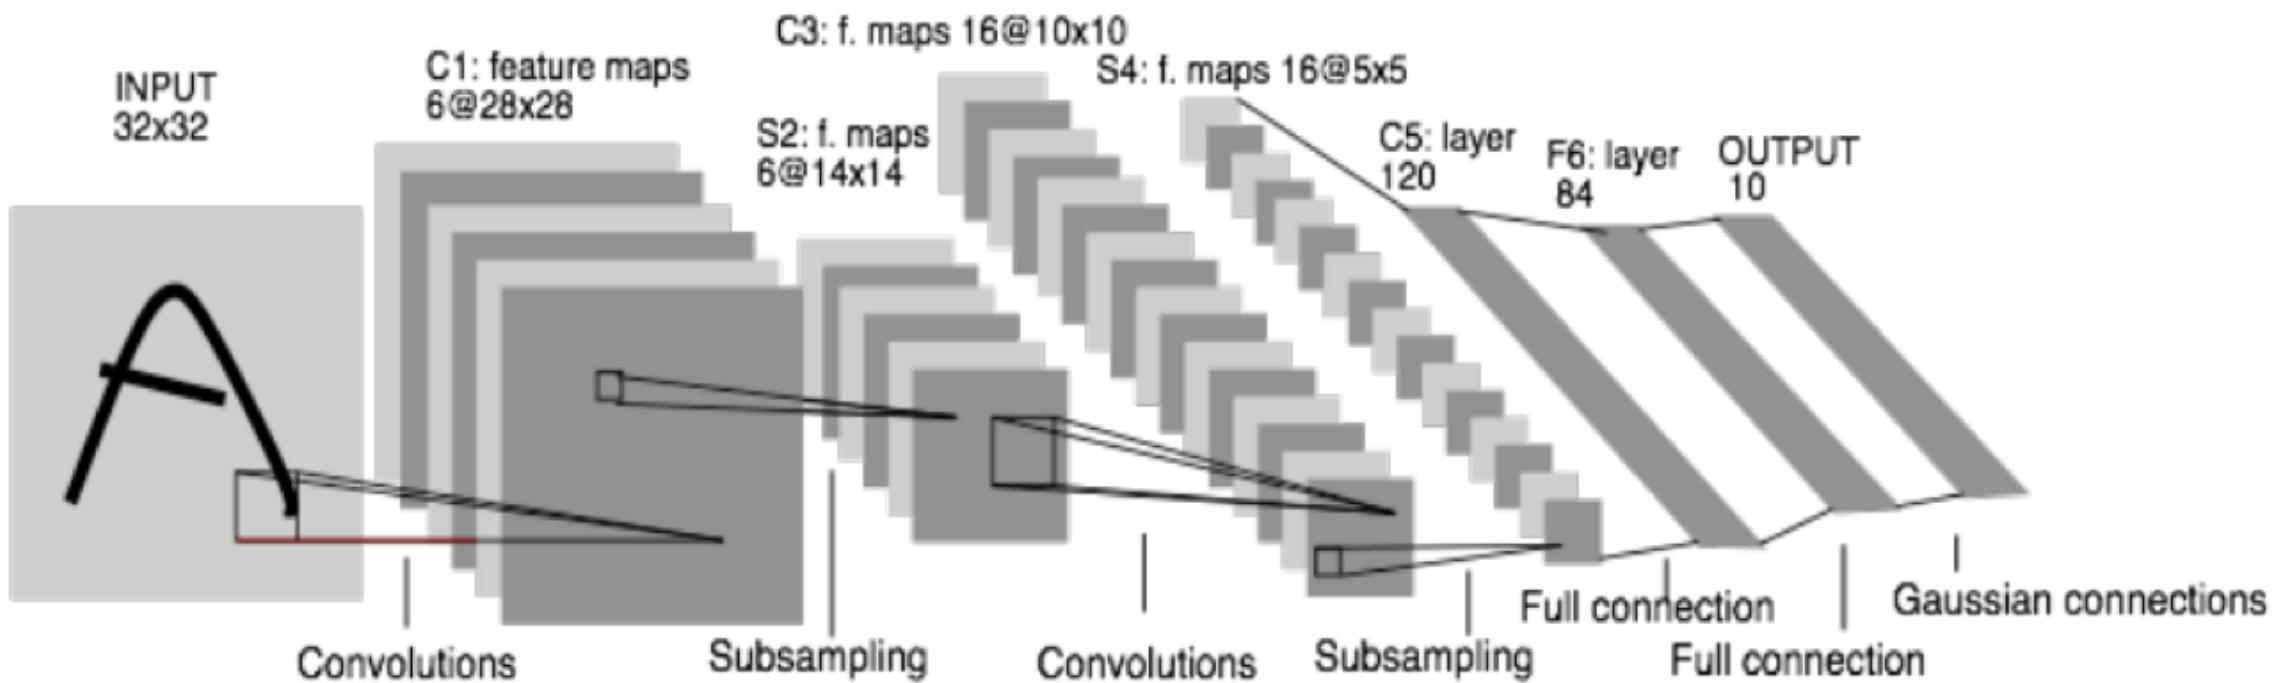
Why Convolutional Neural Networks (CNNs)?

- Goal: reduce number of parameters in fully connected networks; i.e., when each node provides input to each node in the next layer



- Assume 2 layer model with 100 nodes per layer
 - e.g. how many weights are in a 640x480 image?
 - $640 \times 480 \times 3 \times 100 + 100 \times 100 + 100 \times 1 = 92,170,100$
 - e.g. how many weights are in a 2048X1536 image (3.1 Megapixel image)?
 - $2048 \times 1536 \times 3 \times 100 + 100 \times 100 + 100 \times 1 = 943,728,500$

First CNN: LeNet

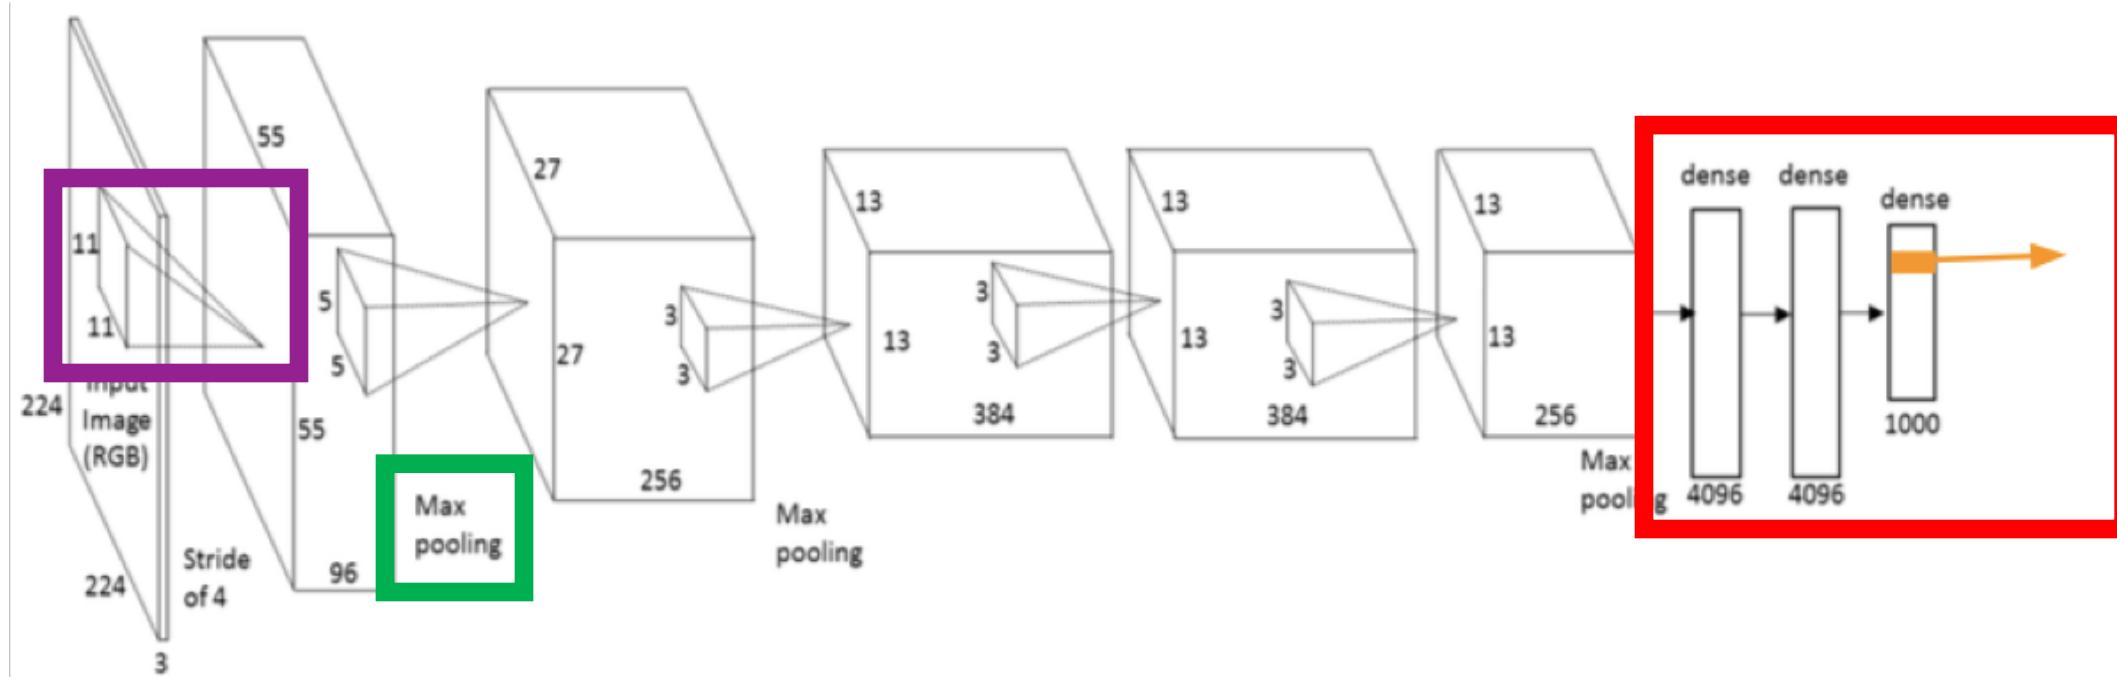


Slide Credit: <https://people.eecs.berkeley.edu/~jrs/189/lec/cnn.pdf>

Y. Lecun ; L. Bottou ; Y. Bengio ; P. Haffner; Gradient-based learning applied to document recognition; 1998

CNN Architecture: Key Components

- CNN extracts useful features of lower dimension prior to passing it to **MLP** with:
 - Convolutional layers
 - Pooling Layers

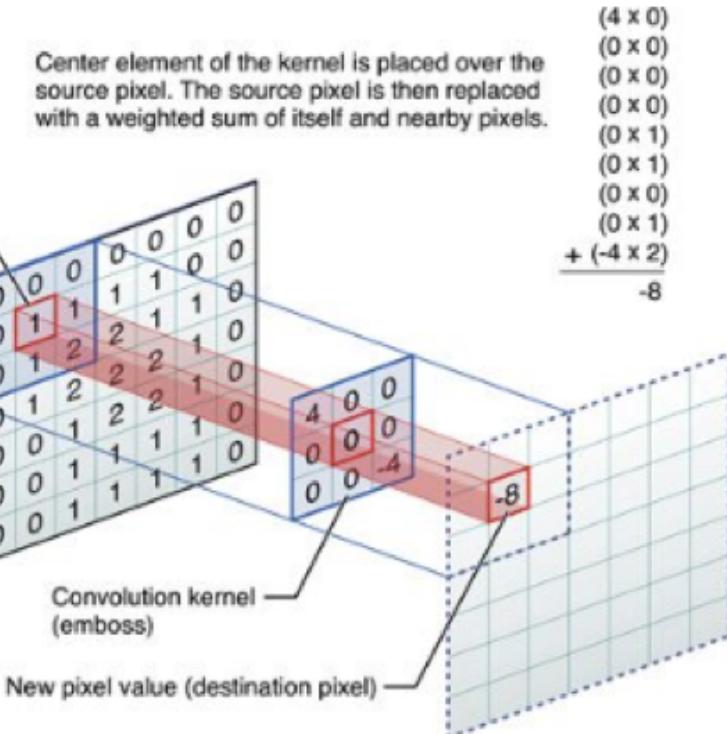
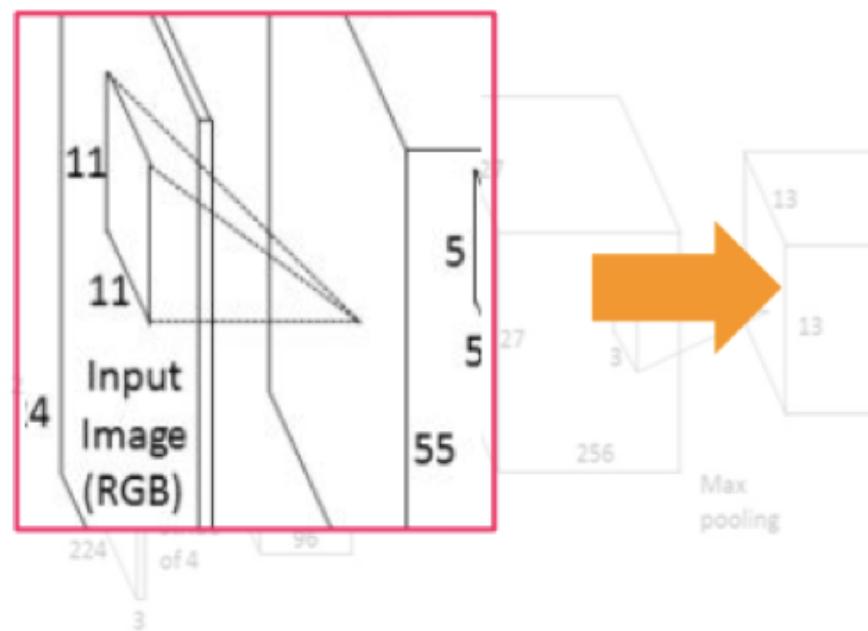


Slide Credit: <https://www.slideshare.net/xavigiro/saliency-prediction-using-deep-learning-techniques>

A. Krizhevsky, I. Sutskever, G. E. Hinton "ImageNet classification with deep convolutional neural networks"

Convolutional Layer

- Set of kernels: each kernel computes the convolution of the output of a previous layer
- Kernels are learned from a given training dataset via *gradient descent*



Convolution Operation: Toy Example

Image				
1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Kernel		
1	0	1
0	1	0
1	0	1

Feature Map		
?	?	?
?	?	?
?	?	?

Convolution Operation: Toy Example

Image				
1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Kernel		
1	0	1
0	1	0
1	0	1

Feature Map		
?	?	?
?	?	?
?	?	?

$$\text{Weighted Sum} = 1 \times 1 + 1 \times 0 + 1 \times 1 + 0 \times 0 + 1 \times 1 + 1 \times 0 + 0 \times 1 + 0 \times 0 + 1 \times 1 = 4$$

Convolution Operation: Toy Example

Image

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Kernel

1	0	1
0	1	0
1	0	1

Feature Map

4	?	?
?	?	?
?	?	?

Convolution Operation: Toy Example

Image				
1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Kernel		
1	0	1
0	1	0
1	0	1

Feature Map		
4	?	?
?	?	?
?	?	?

Weighted Sum = ?

Convolution Operation: Toy Example

Image

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Kernel

1	0	1
0	1	0
1	0	1

Feature Map

4	3	?
?	?	?
?	?	?

Convolution Operation: Toy Example

Image

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Kernel

1	0	1
0	1	0
1	0	1

Feature Map

4	3	?
?	?	?
?	?	?

Weighted Sum = ?

Convolution Operation: Toy Example

Image	Kernel	Feature Map																																											
<table border="1"><tr><td>1</td><td>1</td><td>1</td><td>0</td><td>0</td></tr><tr><td>0</td><td>1</td><td>1</td><td>1</td><td>0</td></tr><tr><td>0</td><td>0</td><td>1</td><td>1</td><td>1</td></tr><tr><td>0</td><td>0</td><td>1</td><td>1</td><td>0</td></tr><tr><td>0</td><td>1</td><td>1</td><td>0</td><td>0</td></tr></table>	1	1	1	0	0	0	1	1	1	0	0	0	1	1	1	0	0	1	1	0	0	1	1	0	0	<table border="1"><tr><td>1</td><td>0</td><td>1</td></tr><tr><td>0</td><td>1</td><td>0</td></tr><tr><td>1</td><td>0</td><td>1</td></tr></table>	1	0	1	0	1	0	1	0	1	<table border="1"><tr><td>4</td><td>3</td><td>4</td></tr><tr><td>?</td><td>?</td><td>?</td></tr><tr><td>?</td><td>?</td><td>?</td></tr></table>	4	3	4	?	?	?	?	?	?
1	1	1	0	0																																									
0	1	1	1	0																																									
0	0	1	1	1																																									
0	0	1	1	0																																									
0	1	1	0	0																																									
1	0	1																																											
0	1	0																																											
1	0	1																																											
4	3	4																																											
?	?	?																																											
?	?	?																																											

Convolution Operation: Toy Example

Image				
1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Kernel		
1	0	1
0	1	0
1	0	1

Feature Map		
4	3	4
?	?	?
?	?	?

Weighted Sum = ?

Convolution Operation: Toy Example

Image

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Kernel

1	0	1
0	1	0
1	0	1

Feature Map

4	3	4
2	?	?
?	?	?

Convolution Operation: Toy Example

Image				
1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Kernel		
1	0	1
0	1	0
1	0	1

Feature Map		
4	3	4
2	?	?
?	?	?

Weighted Sum = ?

Convolution Operation: Toy Example

Image				
1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Kernel		
1	0	1
0	1	0
1	0	1

Feature Map		
4	3	4
2	4	?
?	?	?

Convolution Operation: Toy Example

Image				
1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Kernel		
1	0	1
0	1	0
1	0	1

Feature Map		
4	3	4
2	4	?
?	?	?

Weighted Sum = ?

Convolution Operation: Toy Example

Image	Kernel	Feature Map																																											
<table border="1"><tr><td>1</td><td>1</td><td>1</td><td>0</td><td>0</td></tr><tr><td>0</td><td>1</td><td>1</td><td>1</td><td>0</td></tr><tr><td>0</td><td>0</td><td>1</td><td>1</td><td>1</td></tr><tr><td>0</td><td>0</td><td>1</td><td>1</td><td>0</td></tr><tr><td>0</td><td>1</td><td>1</td><td>0</td><td>0</td></tr></table>	1	1	1	0	0	0	1	1	1	0	0	0	1	1	1	0	0	1	1	0	0	1	1	0	0	<table border="1"><tr><td>1</td><td>0</td><td>1</td></tr><tr><td>0</td><td>1</td><td>0</td></tr><tr><td>1</td><td>0</td><td>1</td></tr></table>	1	0	1	0	1	0	1	0	1	<table border="1"><tr><td>4</td><td>3</td><td>4</td></tr><tr><td>2</td><td>4</td><td>3</td></tr><tr><td>?</td><td>?</td><td>?</td></tr></table>	4	3	4	2	4	3	?	?	?
1	1	1	0	0																																									
0	1	1	1	0																																									
0	0	1	1	1																																									
0	0	1	1	0																																									
0	1	1	0	0																																									
1	0	1																																											
0	1	0																																											
1	0	1																																											
4	3	4																																											
2	4	3																																											
?	?	?																																											

Convolution Operation: Toy Example

Image				
1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Kernel		
1	0	1
0	1	0
1	0	1

Feature Map		
4	3	4
2	4	3
?	?	?

Weighted Sum = ?

Convolution Operation: Toy Example

Image				
1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Kernel		
1	0	1
0	1	0
1	0	1

Feature Map		
4	3	4
2	4	3
2	?	?

Convolution Operation: Toy Example

Image				
1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Kernel		
1	0	1
0	1	0
1	0	1

Feature Map		
4	3	4
2	4	3
2	?	?

Weighted Sum = ?

Convolution Operation: Toy Example

Image				
1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Kernel		
1	0	1
0	1	0
1	0	1

Feature Map		
4	3	4
2	4	3
2	3	?

Convolution Operation: Toy Example

Image				
1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Kernel		
1	0	1
0	1	0
1	0	1

Feature Map		
4	3	4
2	4	3
2	3	?

Weighted Sum = ?

Convolution Operation: Toy Example

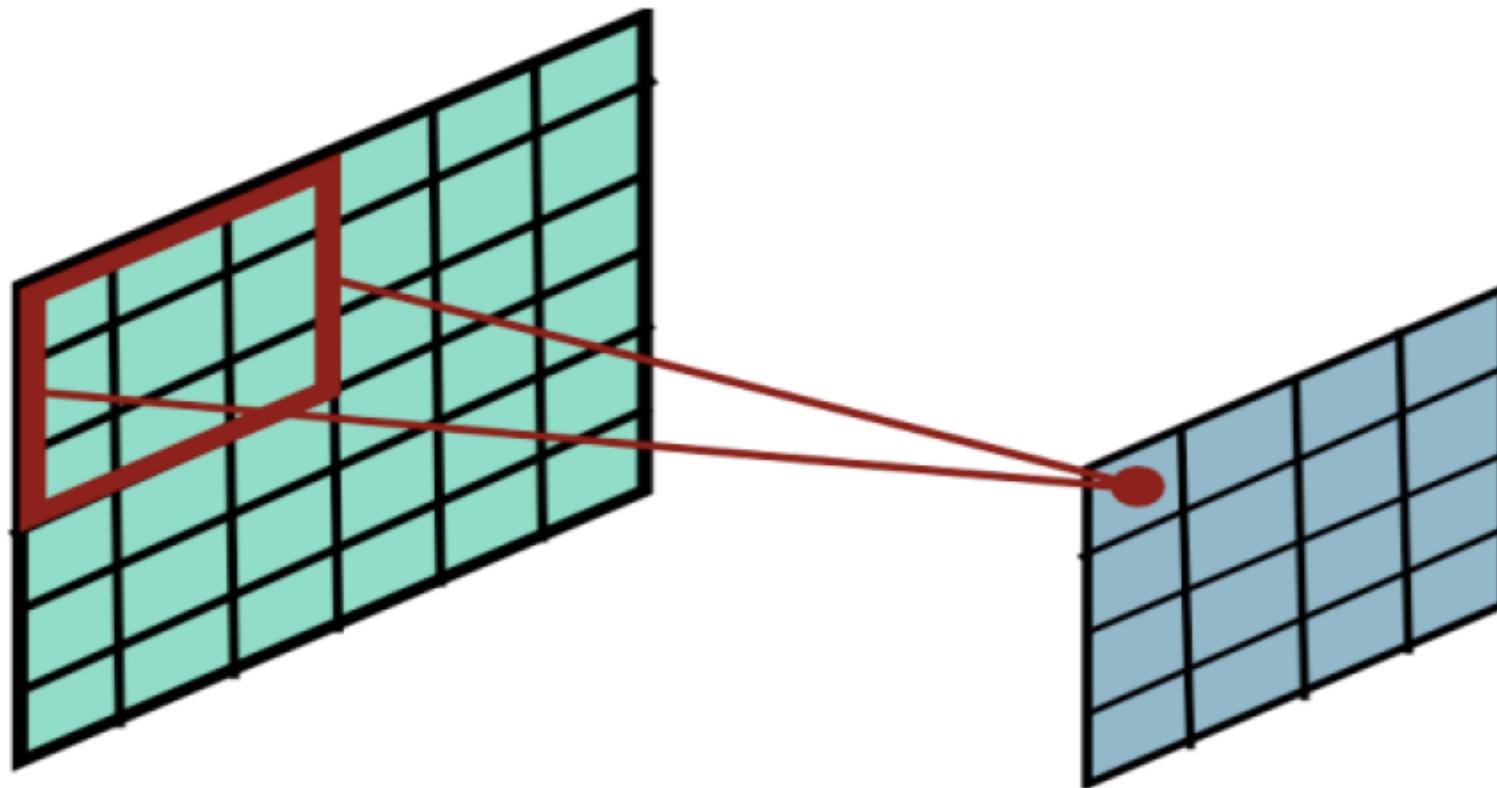
Image				
1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Kernel		
1	0	1
0	1	0
1	0	1

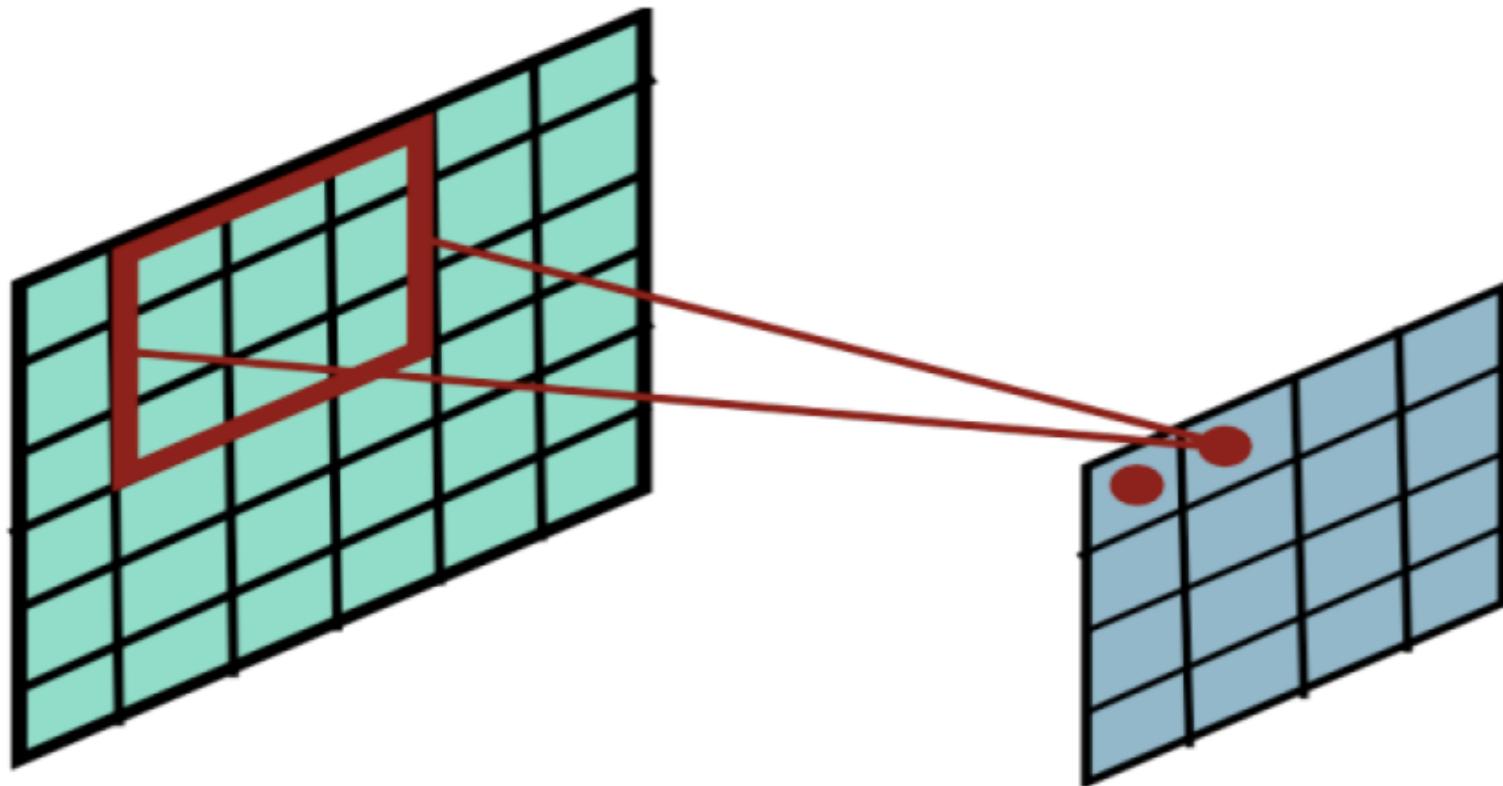
Feature Map		
4	3	4
2	4	3
2	3	4

Weighted Sum = ?

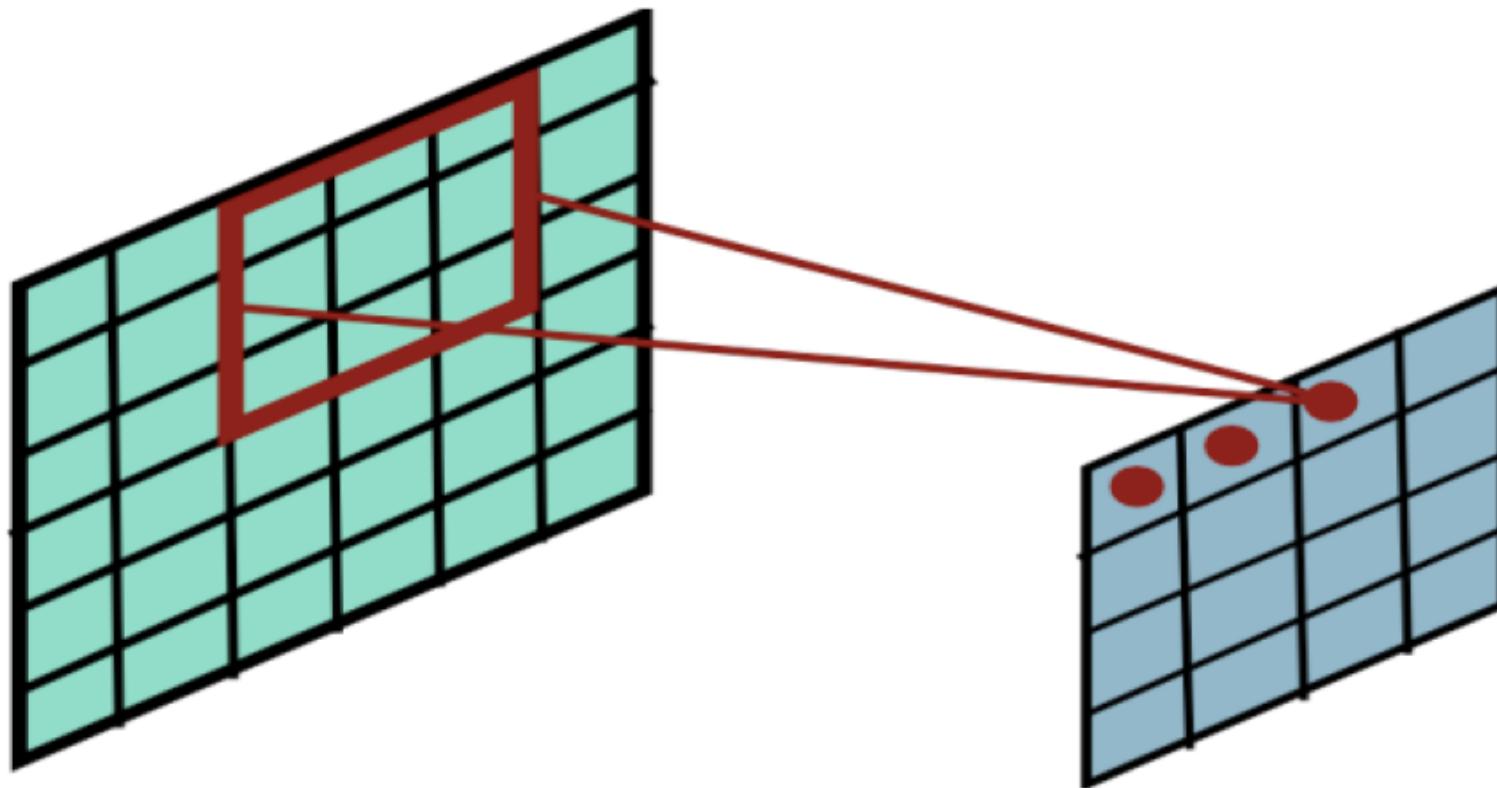
Convolutional Operation (3x3 kernel)



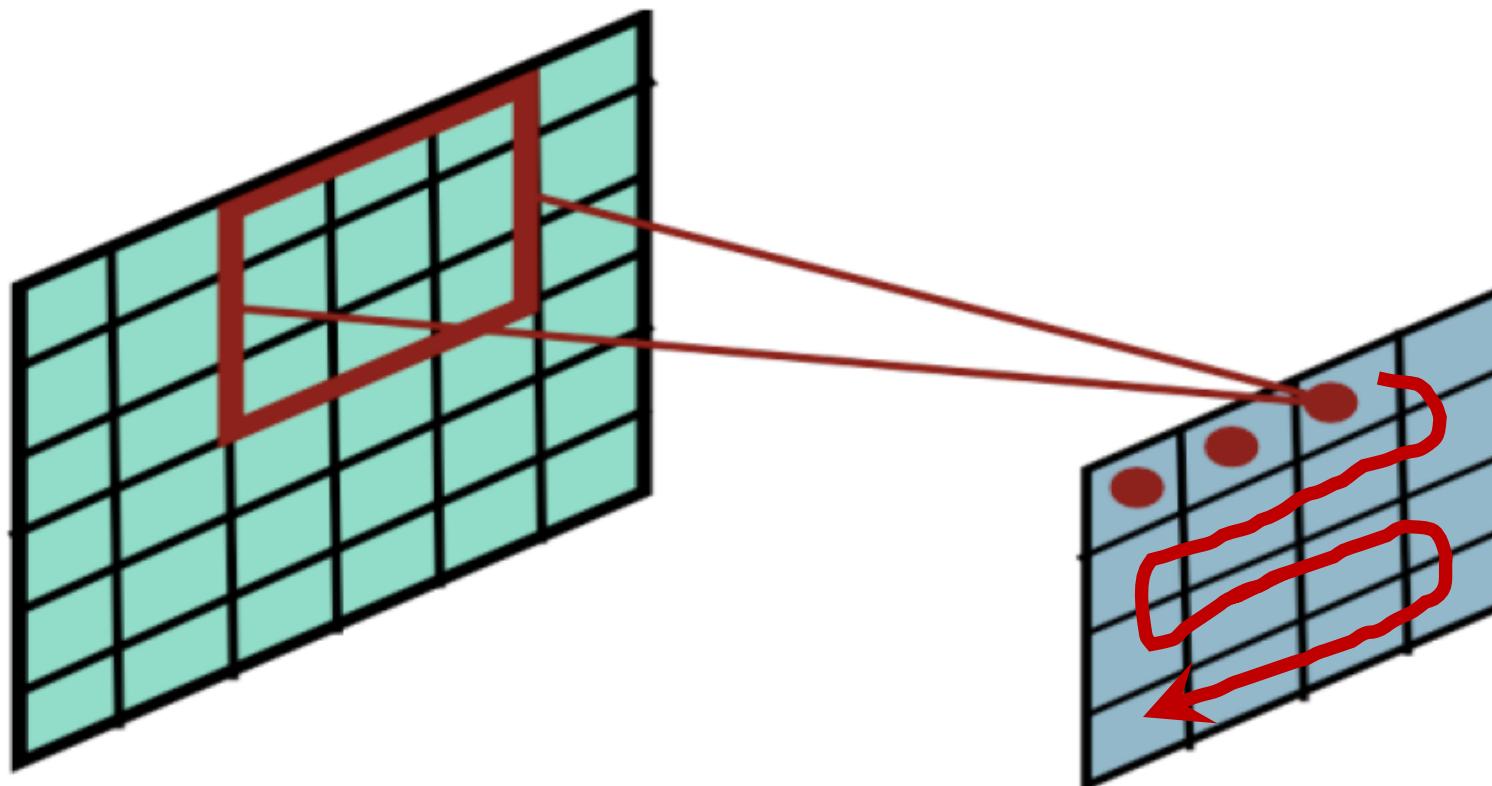
Convolutional Operation (3x3 kernel)



Convolutional Operation (3x3 kernel)



Convolutional Operation (3x3 kernel)



Convolution: What Does It Do?

- e.g.,

Filter

0	0	0	0	0	30	0
0	0	0	0	30	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	0	0	0	0

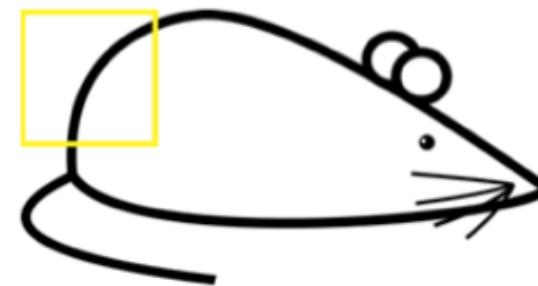
Visualization of Filter



Convolution: What Does It Do?

- e.g.,

Filter Overlaid on Image



Image

0	0	0	0	0	0	30
0	0	0	0	50	50	50
0	0	0	20	50	0	0
0	0	0	50	50	0	0
0	0	0	50	50	0	0
0	0	0	50	50	0	0
0	0	0	50	50	0	0

*

Filter

0	0	0	0	0	30	0
0	0	0	0	30	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	0	0	0	0

Weighted Sum = ?

Weighted Sum = $(50 \times 30) + (20 \times 30) + (50 \times 30) + (50 \times 3) + (50 \times 30)$

Weighted Sum = 6600 (**Large Number!!**)

Convolution: What Does It Do?

- e.g.,

Filter Overlaid on Image



Image

0	0	0	0	0	0	0
0	40	0	0	0	0	0
40	0	40	0	0	0	0
40	20	0	0	0	0	0
0	50	0	0	0	0	0
0	0	50	0	0	0	0
25	25	0	50	0	0	0

*

Filter

0	0	0	0	0	30	0
0	0	0	0	30	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	0	0	0	0

Weighted Sum = ?

Weighted Sum = 0 (**Small Number!!**)

Convolution: What Does It Do?

This Filter is a Curve Detector!

- e.g.,

0	0	0	0	0	30	0
0	0	0	0	30	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	0	0	0	0



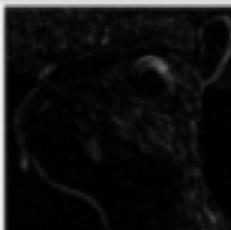
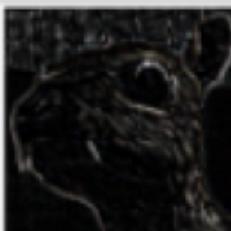
Filter Overlaid on Image (Big Response!)

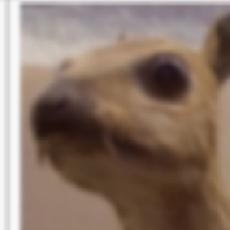


Filter Overlaid on Image (Small Response!)



Different Filters Detect Different Features

	Filter	Feature Map
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	

	Filter	Feature Map
Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	
Box blur (normalized)	$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	
Gaussian blur (approximation)	$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$	

<https://ujjwalkarn.me/2016/08/11/intuitive-explanation-convnets/>

http://www1.adept.com/main/KE/DATA/ACE/AdeptSight_User/ImageProcessing_Operations.html

Convolution Implementation Details

- **Kernel size:** e.g., 3x3, 5x5, 7x7 averaging filters

1	1	1
1	1	1
1	1	1

1	1	1	1	1
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1

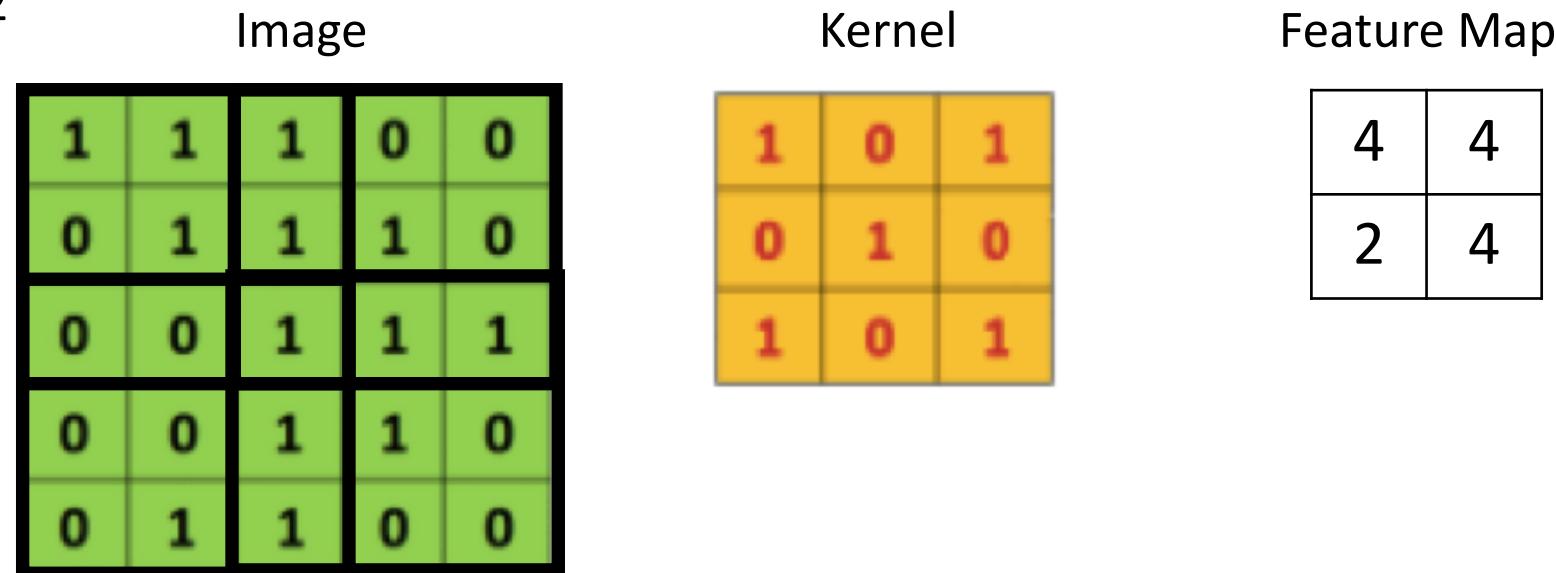
1	1	1	1	1	1	1
1	1	1	1	1	1	1
1	1	1	1	1	1	1
1	1	1	1	1	1	1
1	1	1	1	1	1	1
1	1	1	1	1	1	1
1	1	1	1	1	1	1

- Why choose a larger versus smaller kernel?

Convolution Implementation Details

- **Stride:** how many steps taken spatially before applying a filter

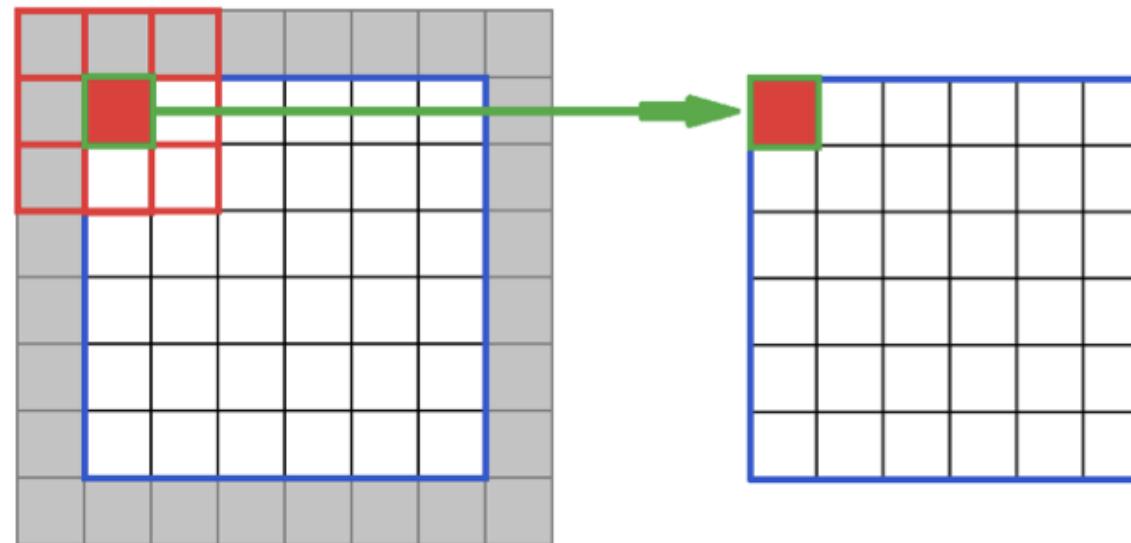
- e.g., 2x2



- Why choose a larger versus smaller stride size?

Convolution Implementation Details

- **Padding:** add zeros at the image boundaries
 - e.g., “same” padding versus “valid” (no) padding



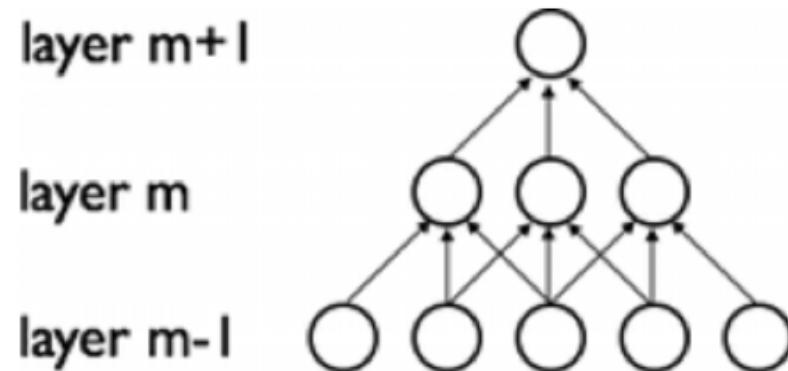
- Why use padding? How much padding to use?

Demo

http://deeplearning.net/software/theano/tutorial/conv_arithmetic.html

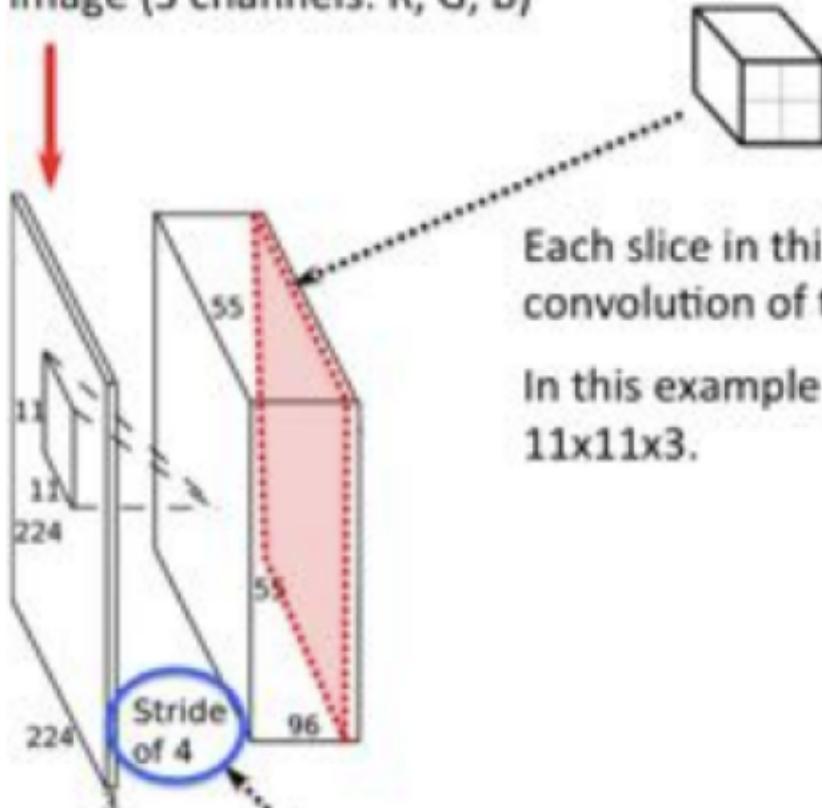
Convolutional Layer

- Implementation:
 1. Apply convolution operation with each filter
 2. Add biases (one per each output image).
 3. Apply an activation function to all the pixels of the output images.
- CNNs learn “filters” based on spatially **local input patterns**.
- Stacking many convolutional layers leads to learning patterns in increasingly **larger regions of the pixel space**



Recall CNN Architecture

image (3 channels: R, G, B)



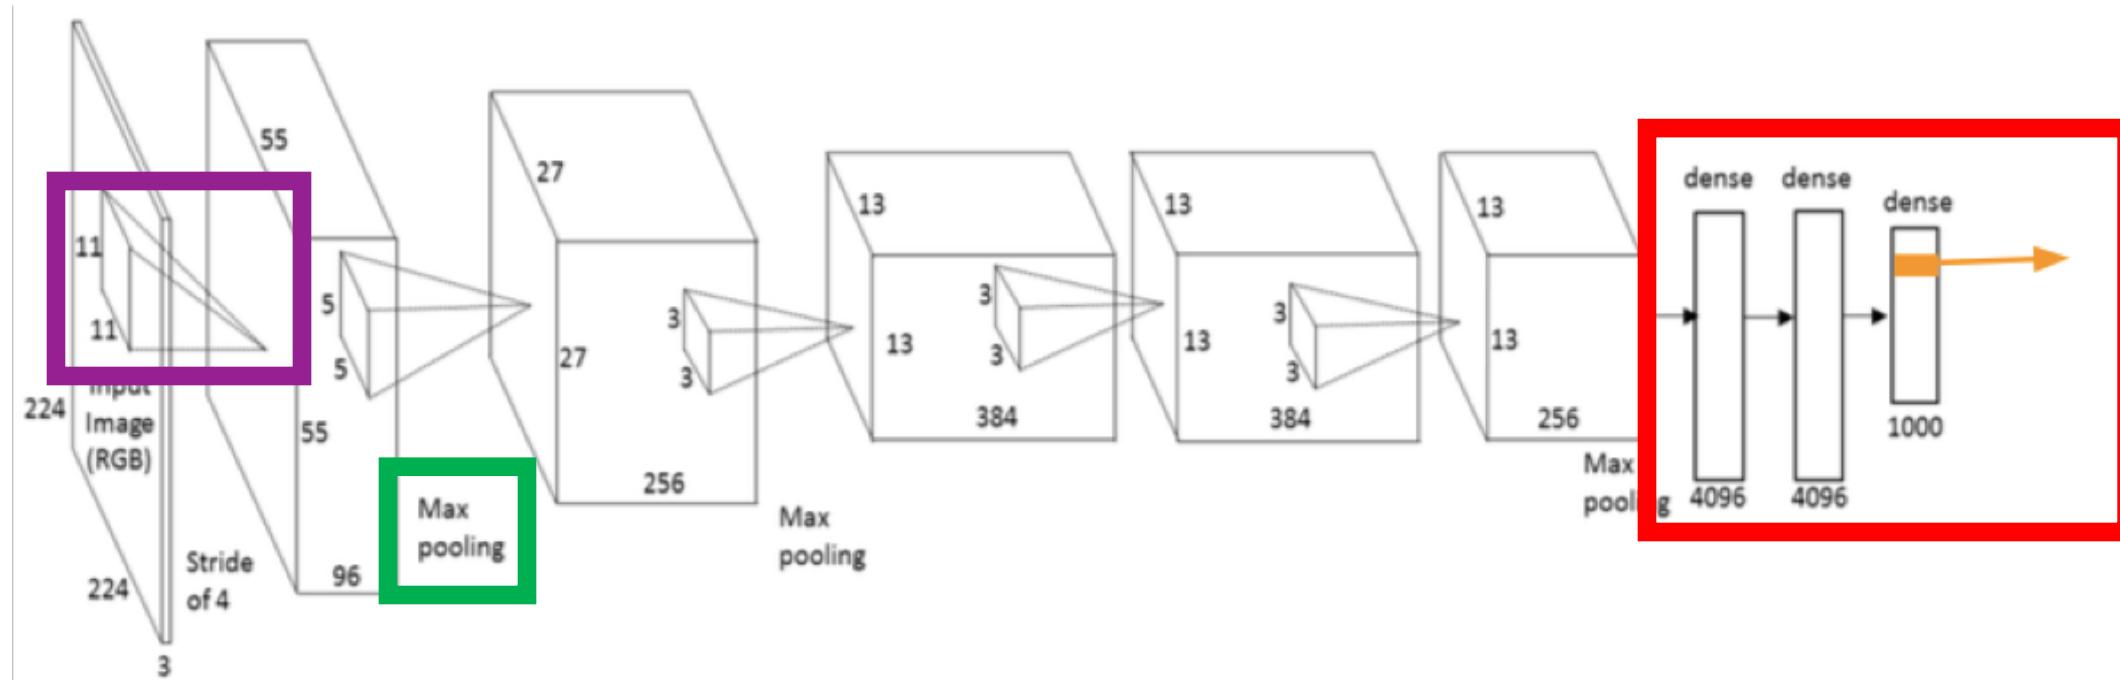
Each slice in this cube is the output of convolution of the image and a filter.

In this example the filter size is
11x11x3.

We don't do convolution in every pixel, but in
every 4th pixel (in x and y direction)

Recall CNN Architecture: Key Components

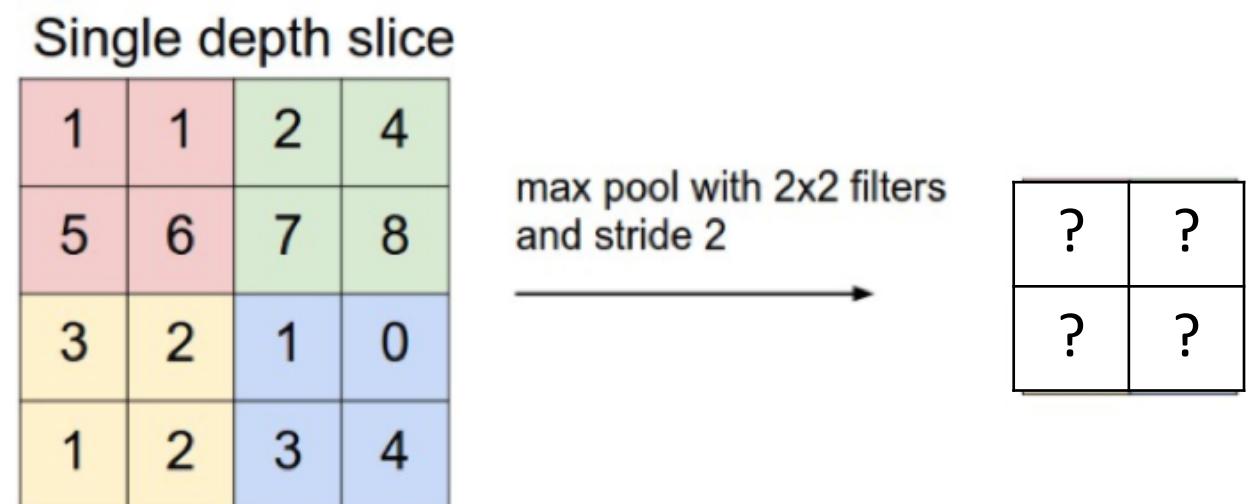
- CNN extracts useful features of lower dimension prior to passing it to **MLP** with:
 - Convolutional layers
 - Pooling Layers



Slide Credit: <https://www.slideshare.net/xavigiro/saliency-prediction-using-deep-learning-techniques>
A. Krizhevsky, I. Sutskever, G. E. Hinton "ImageNet classification with deep convolutional neural networks"

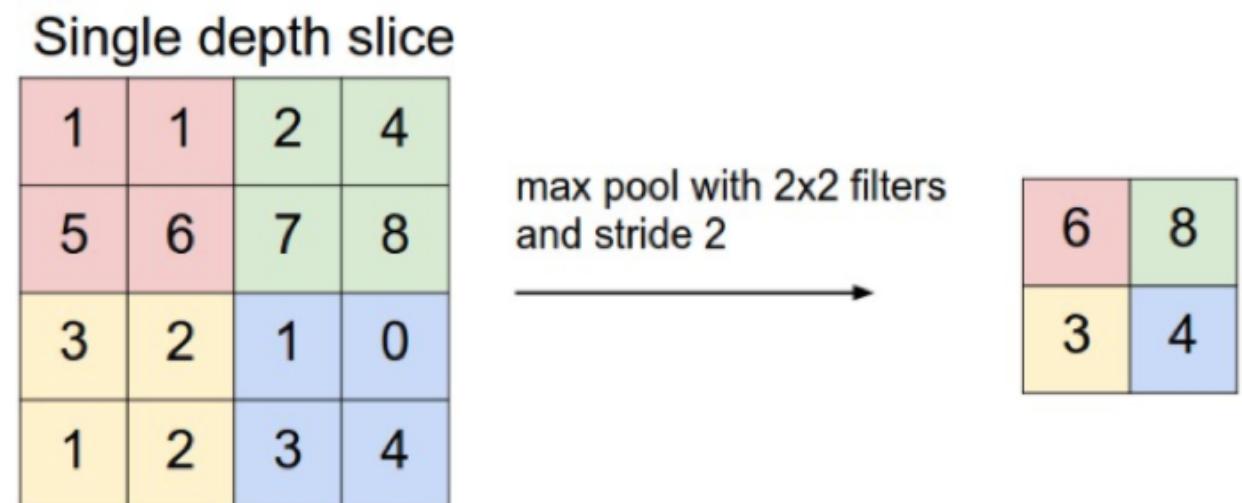
Pooling Layer

- Reduces feature dimension size
 - **Max-pooling:** partitions input into a set of non-overlapping rectangles and outputs the maximum value for each chunk



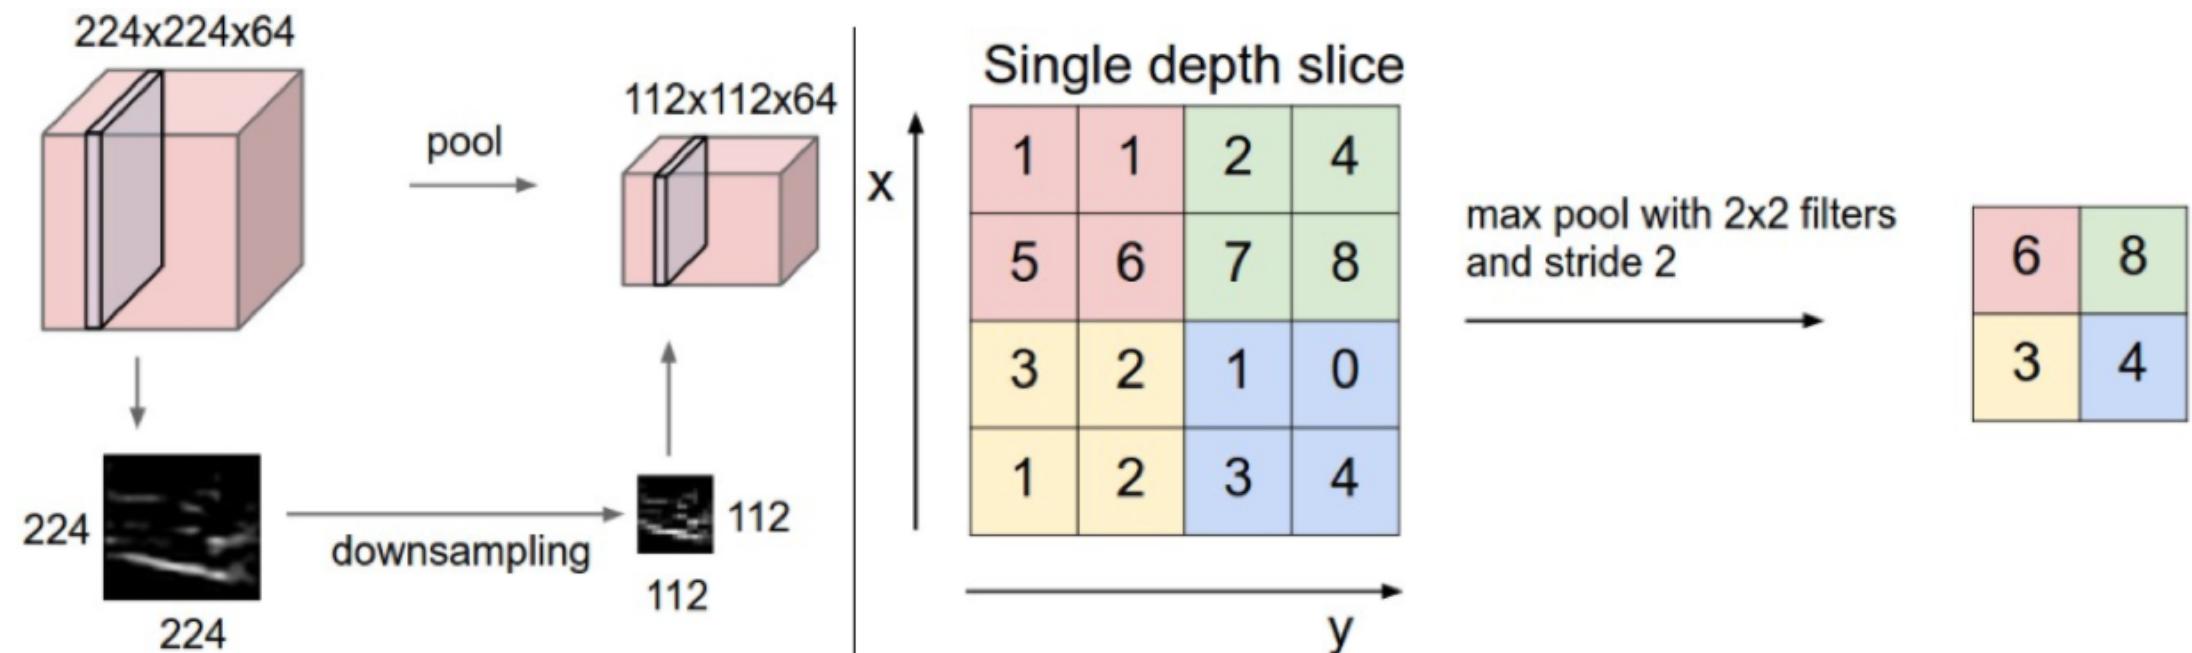
Pooling Layer

- Reduces feature dimension size
 - **Max-pooling:** partitions input into a set of non-overlapping rectangles and outputs the maximum value for each chunk



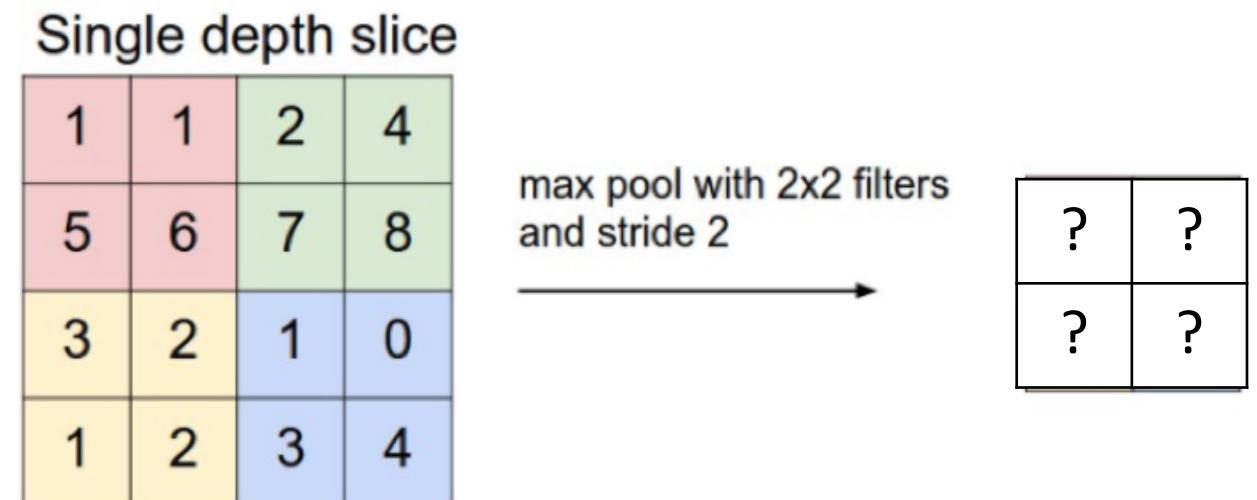
Pooling Layer

- Reduces feature dimension size
 - **Max-pooling:** partitions input into a set of non-overlapping rectangles and outputs the maximum value for each chunk



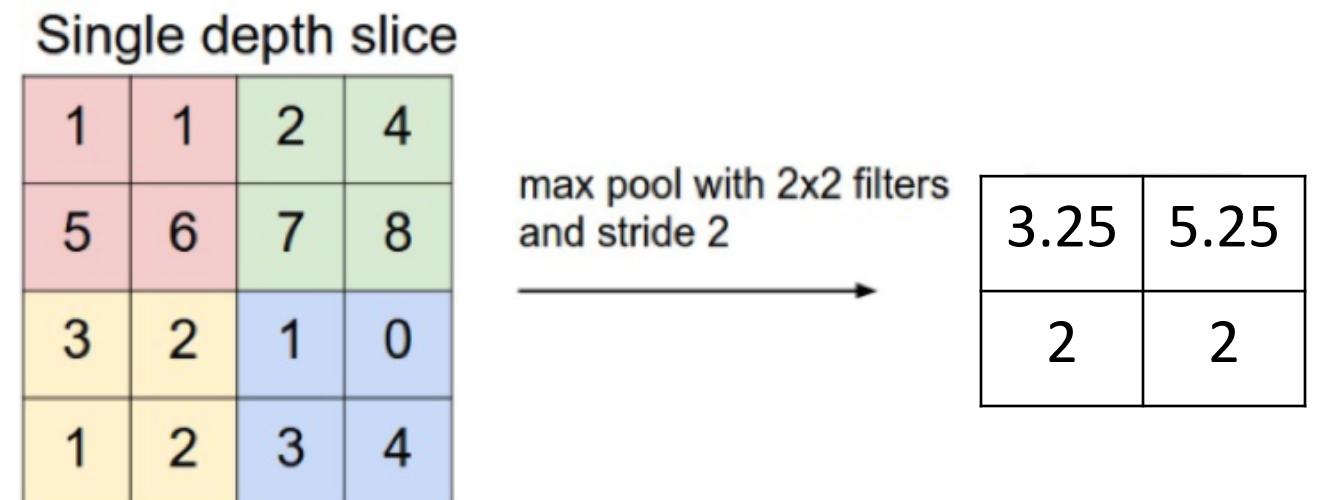
Pooling Layer

- Reduces feature dimension size
 - **Max-pooling:** partitions input into a set of non-overlapping rectangles and outputs the maximum value for each chunk
 - **Average-pooling:** partitions input into a set of non-overlapping rectangles and outputs the average value for each chunk

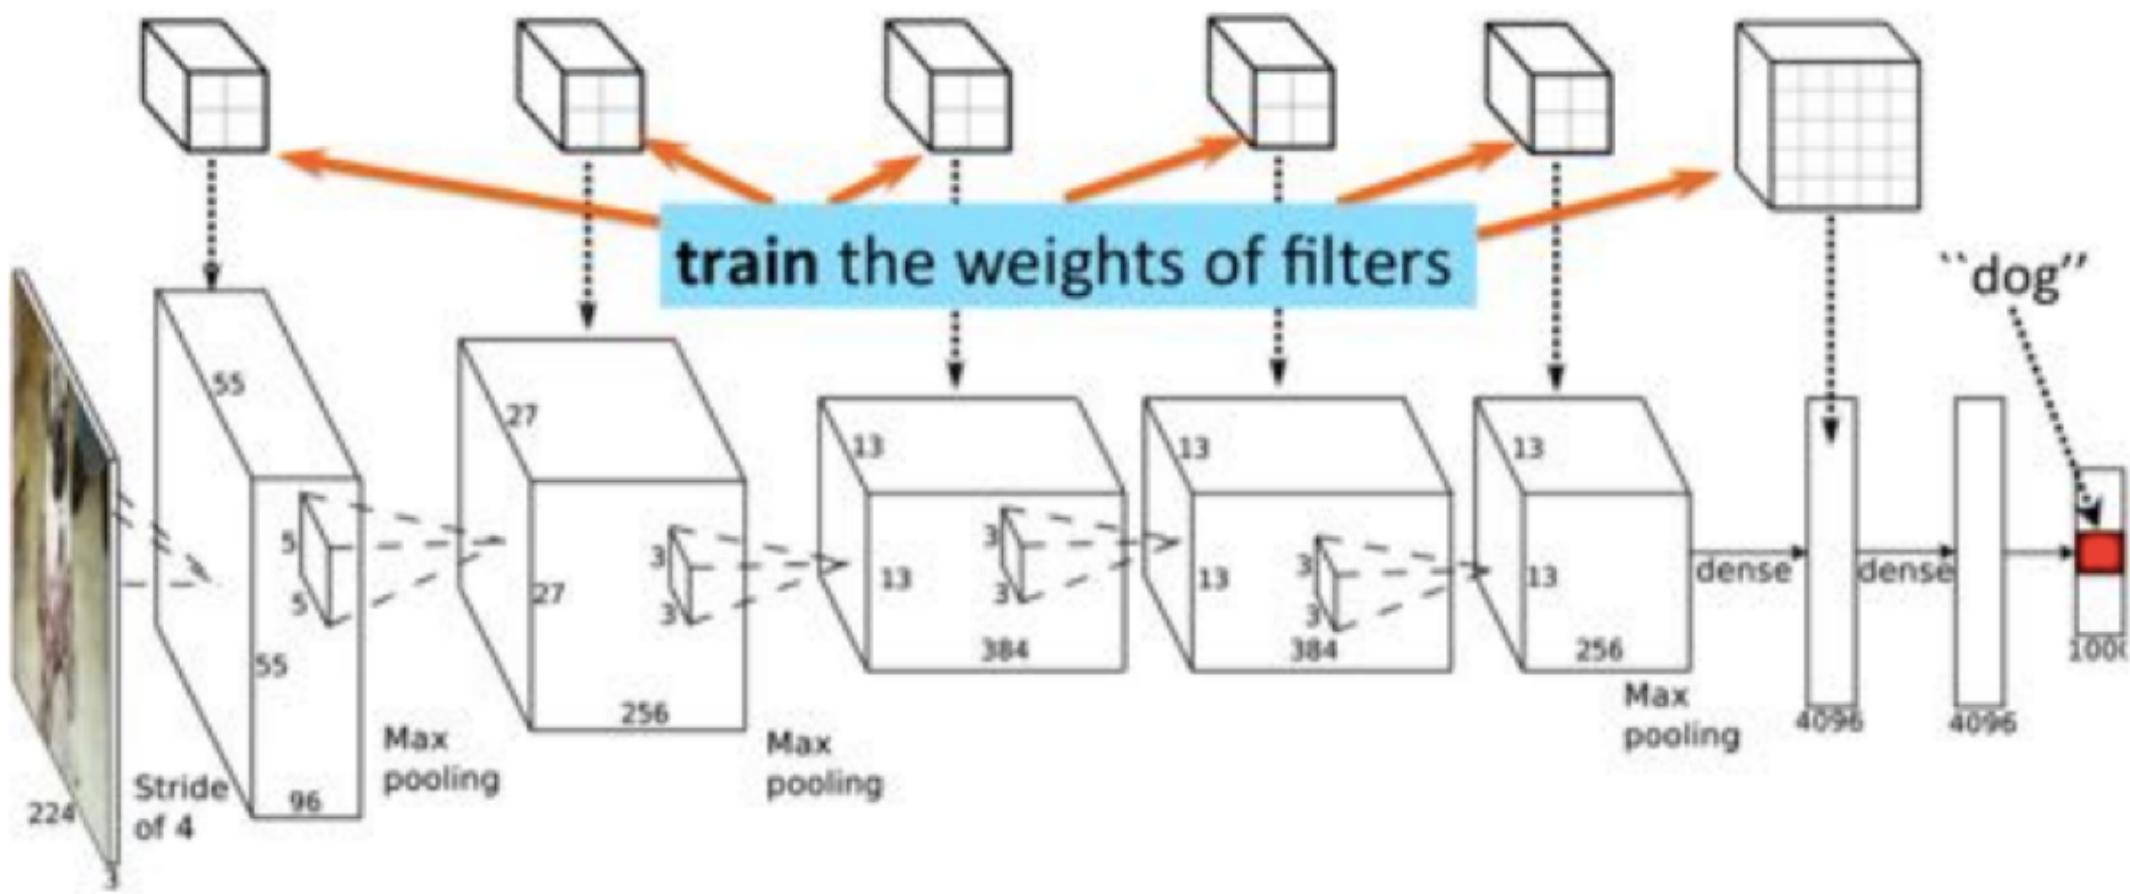


Pooling Layer

- Reduces feature dimension size
 - **Max-pooling:** partitions input into a set of non-overlapping rectangles and outputs the maximum value for each chunk
 - **Average-pooling:** partitions input into a set of non-overlapping rectangles and outputs the average value for each chunk

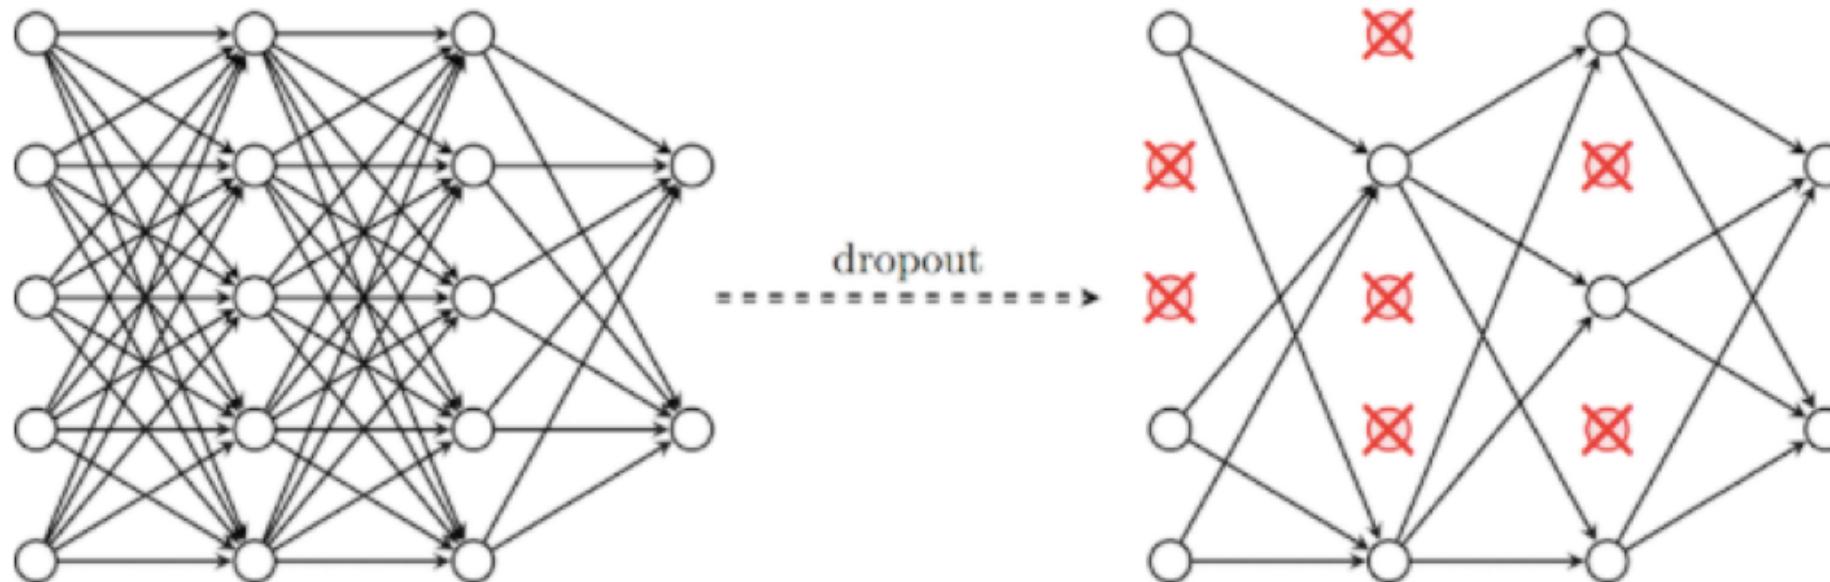


Training CNN

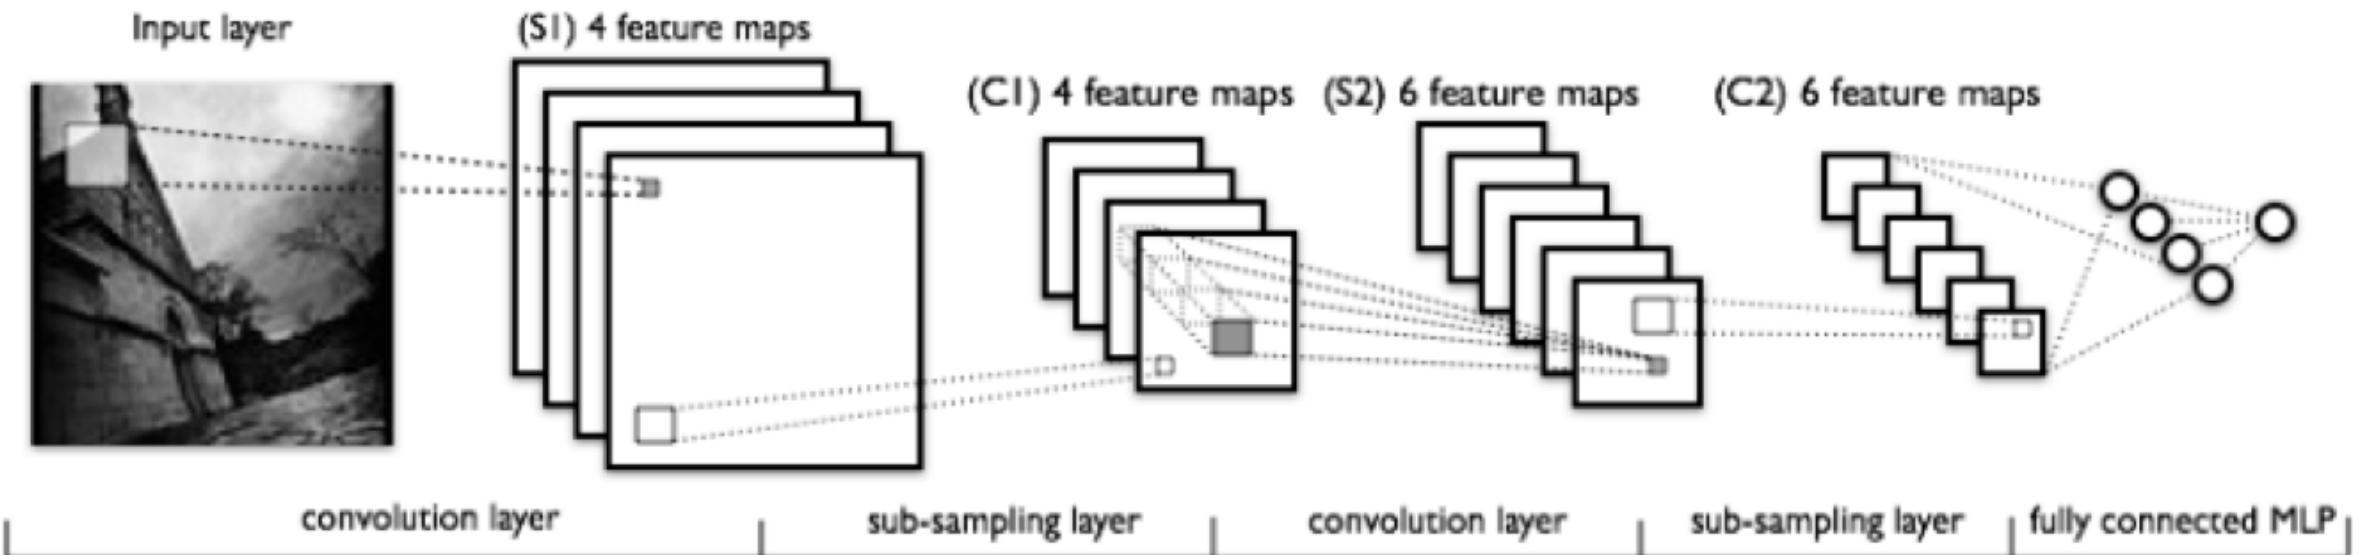


Training CNN: Dropout

- In single training iteration, *eliminate nodes from the network*
- Why? This forces the neural network to learn redundancy by developing a consensus (ensemble) of neurons within a layer that are predictive



Revisiting LeNet

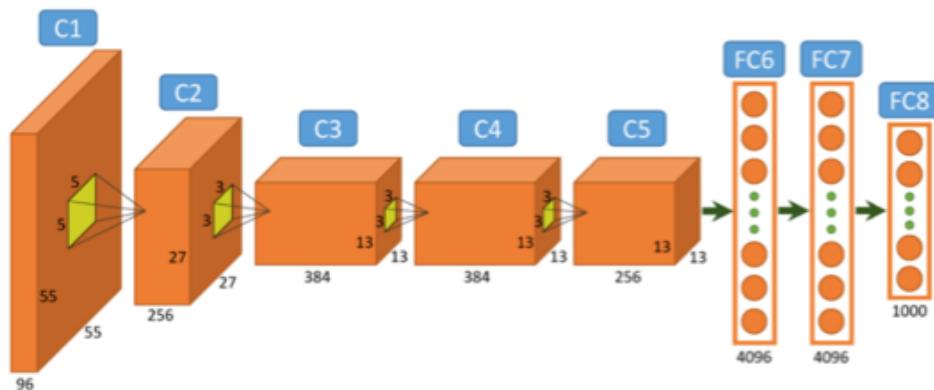


<http://deeplearning.net/tutorial/lenet.html>

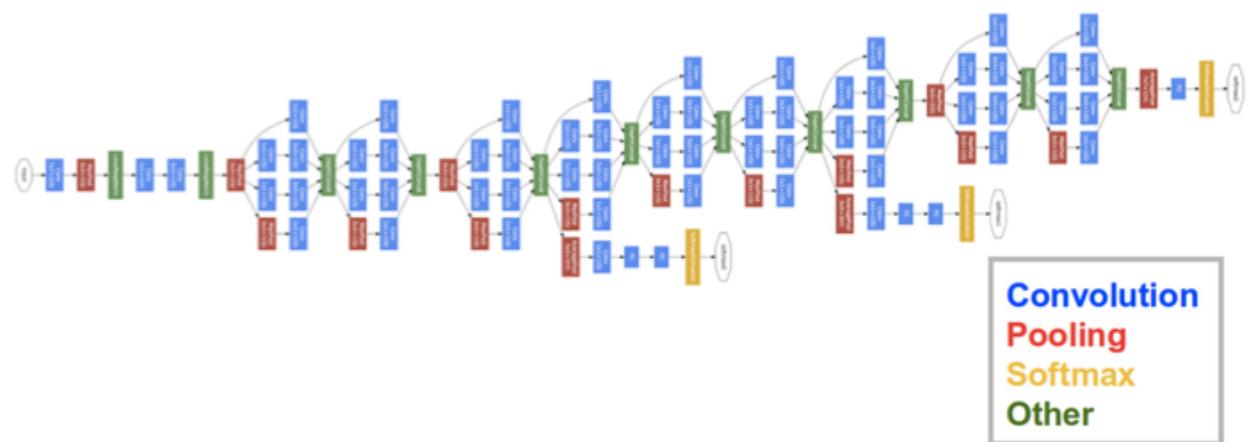
Y. Lecun ; L. Bottou ; Y. Bengio ; P. Haffner; Gradient-based learning applied to document recognition; 1998

Over Time, CNNs are Becoming “Deeper”

AlexNet (2012)



GoogLeNet (2014)

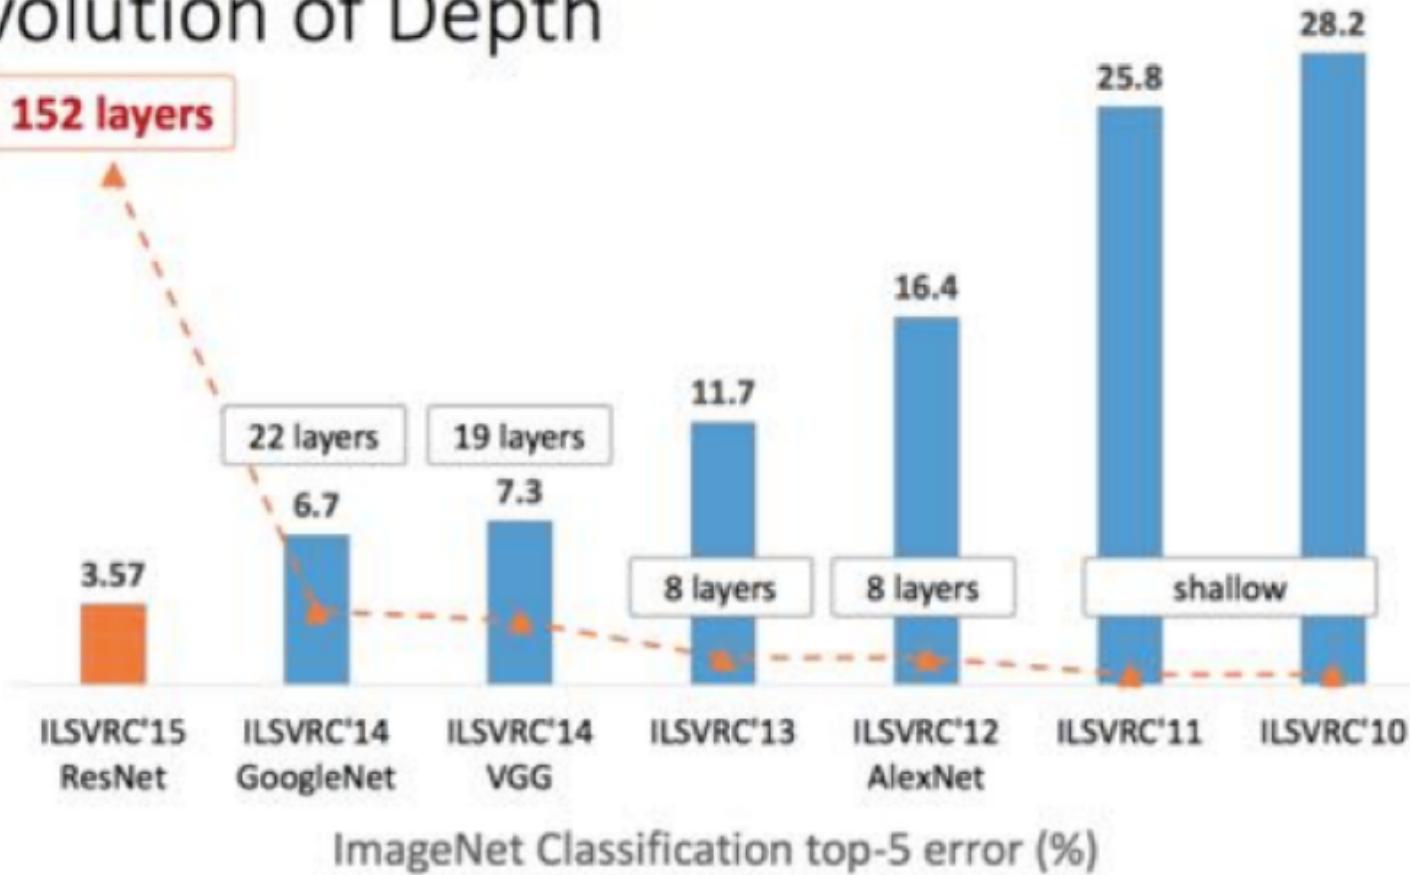


http://joelouismarino.github.io/blog_posts/blog_googlenet_keras.html

https://www.researchgate.net/figure/Architecture-of-Alexnet-From-left-to-right-input-to-output-five-convolutional-layers_fig2_312303454

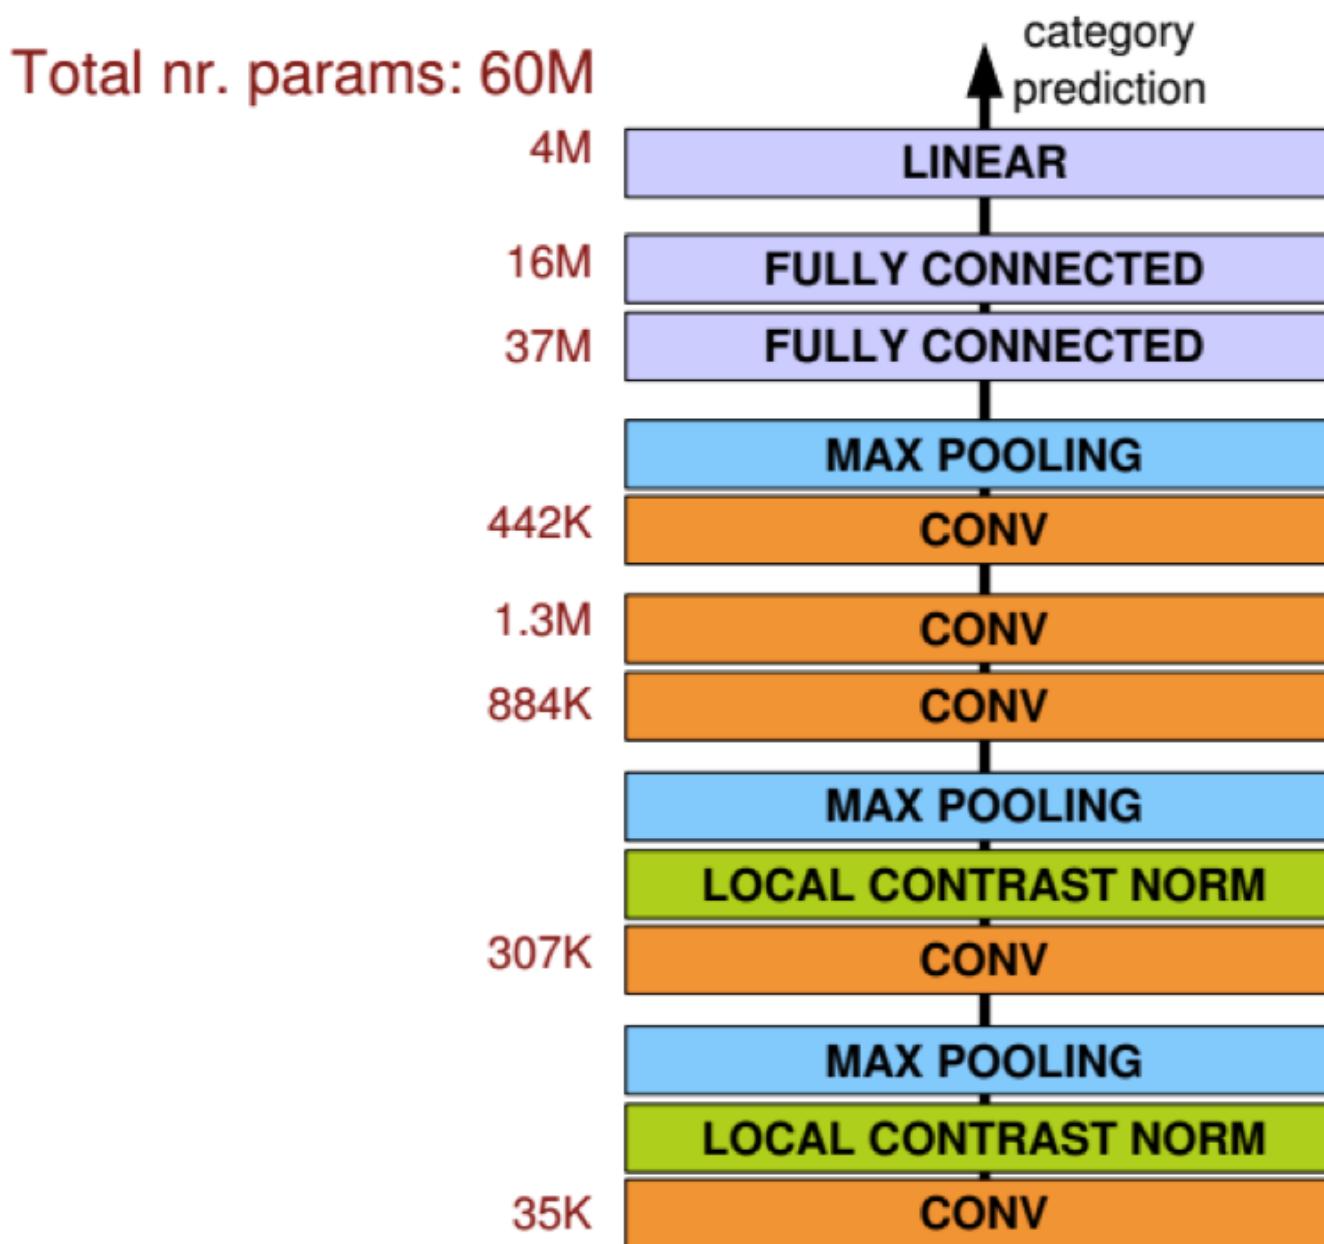
Over Time, CNNs are Becoming “Deeper”

Revolution of Depth



Slide: R. Liao, Paper: [He, K., Zhang, X., Ren, S. and Sun, J., 2015. Deep Residual Learning for Image Recognition. arXiv:1512.03385, 2016]

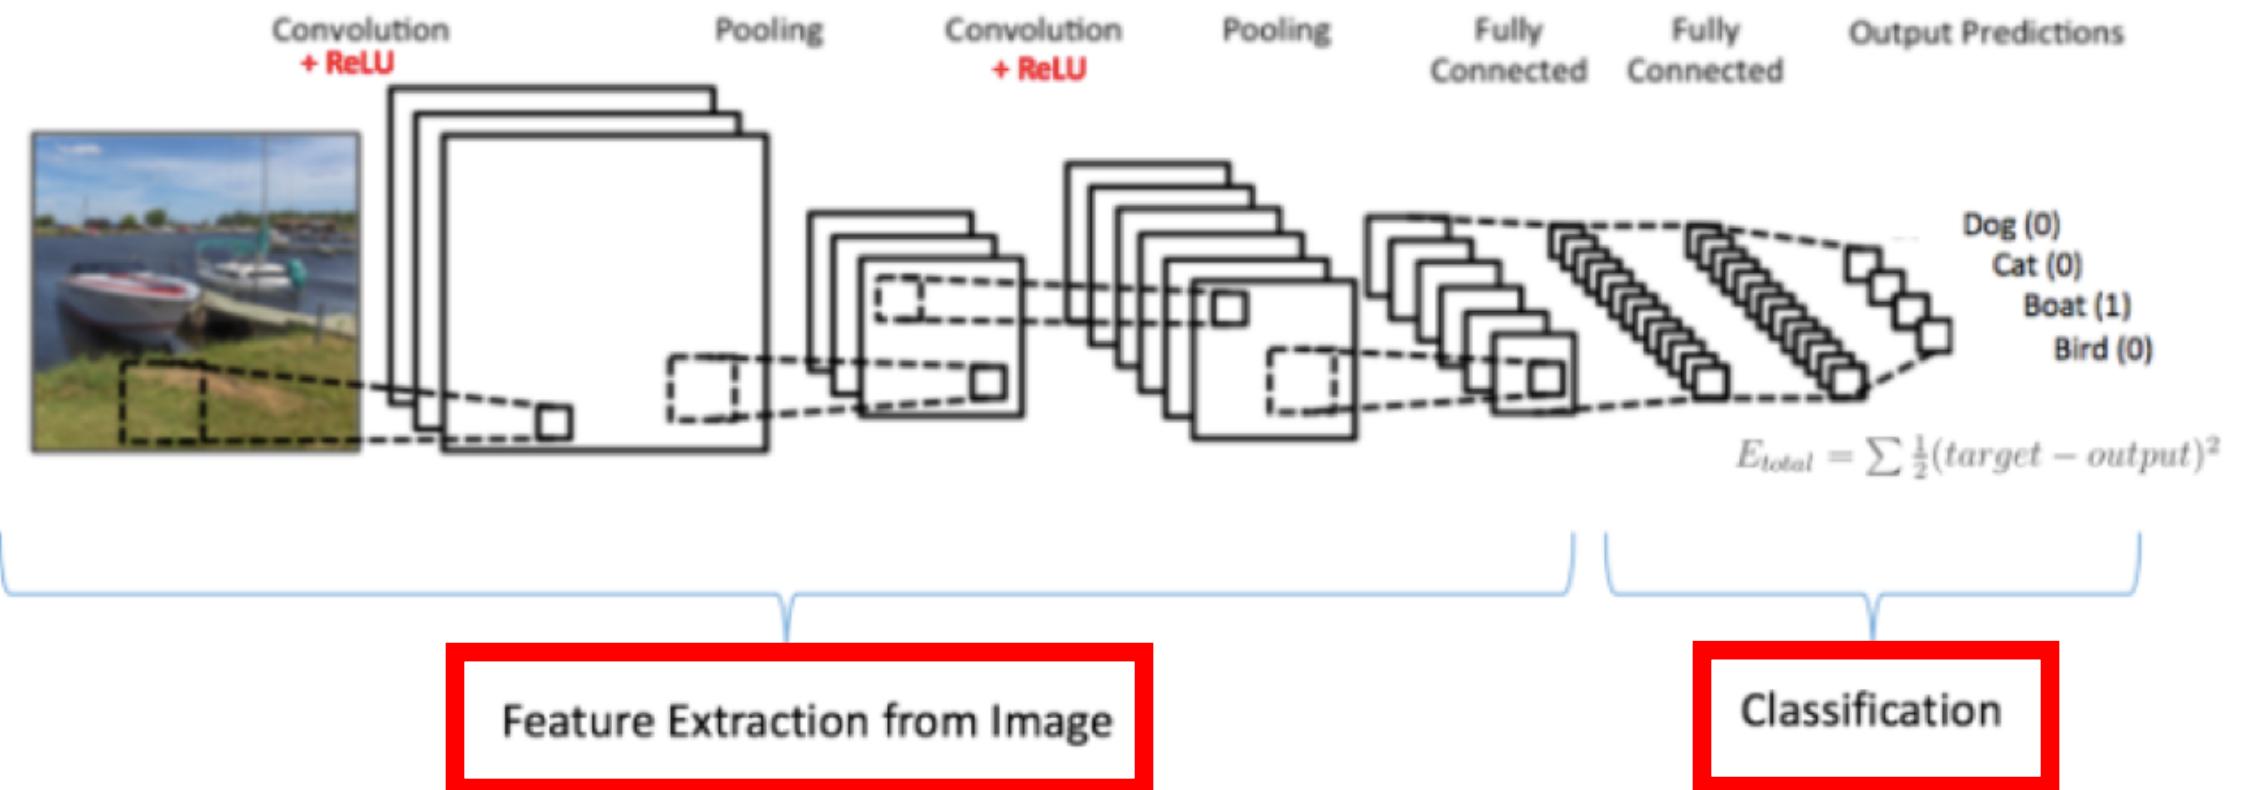
AlexNet



Today's Topics

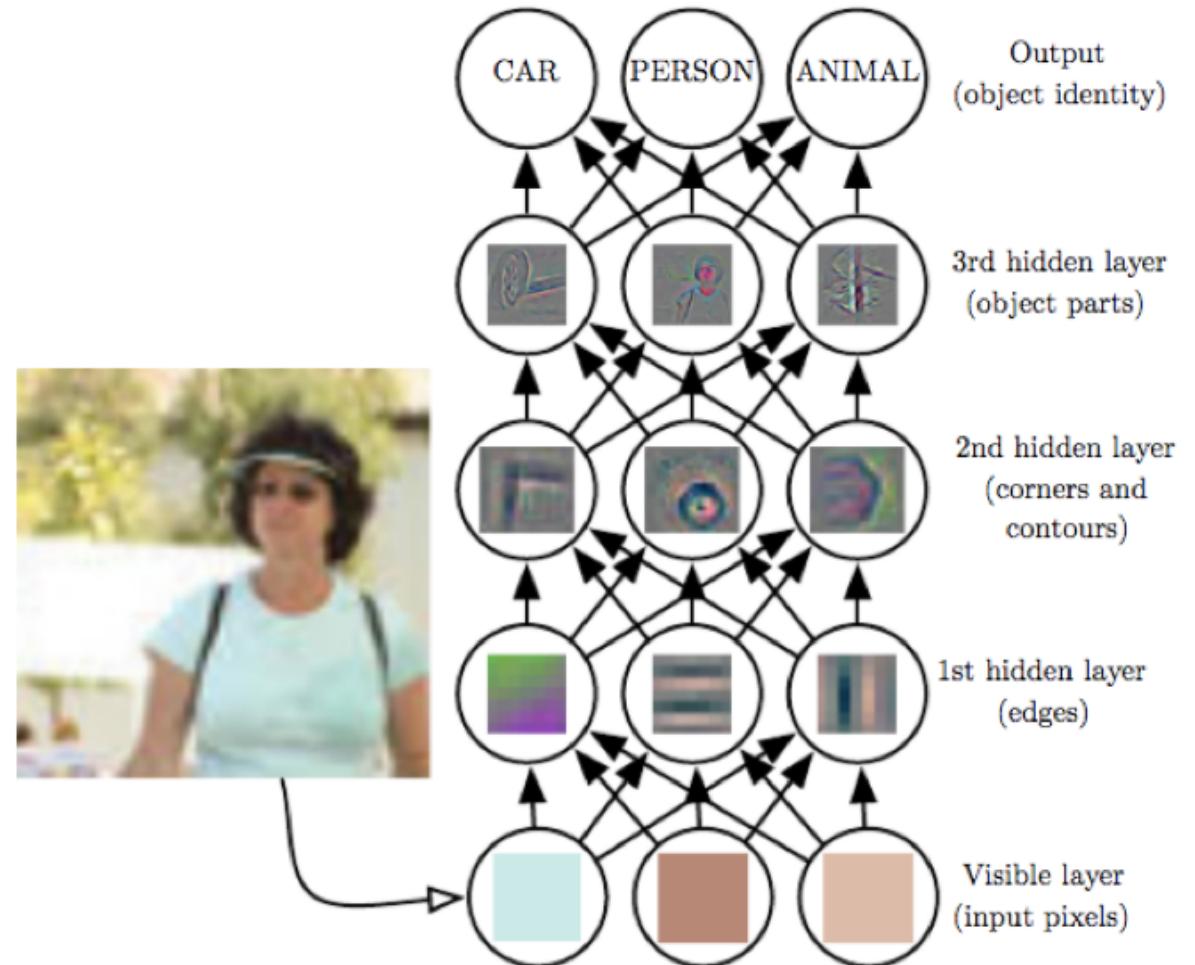
- Revisiting History of Modern Neural Networks and “Deep Learning”
- Convolutional Neural Networks (CNN)
- Deep Features
- Lab

CNN: Intuition of Different Layers



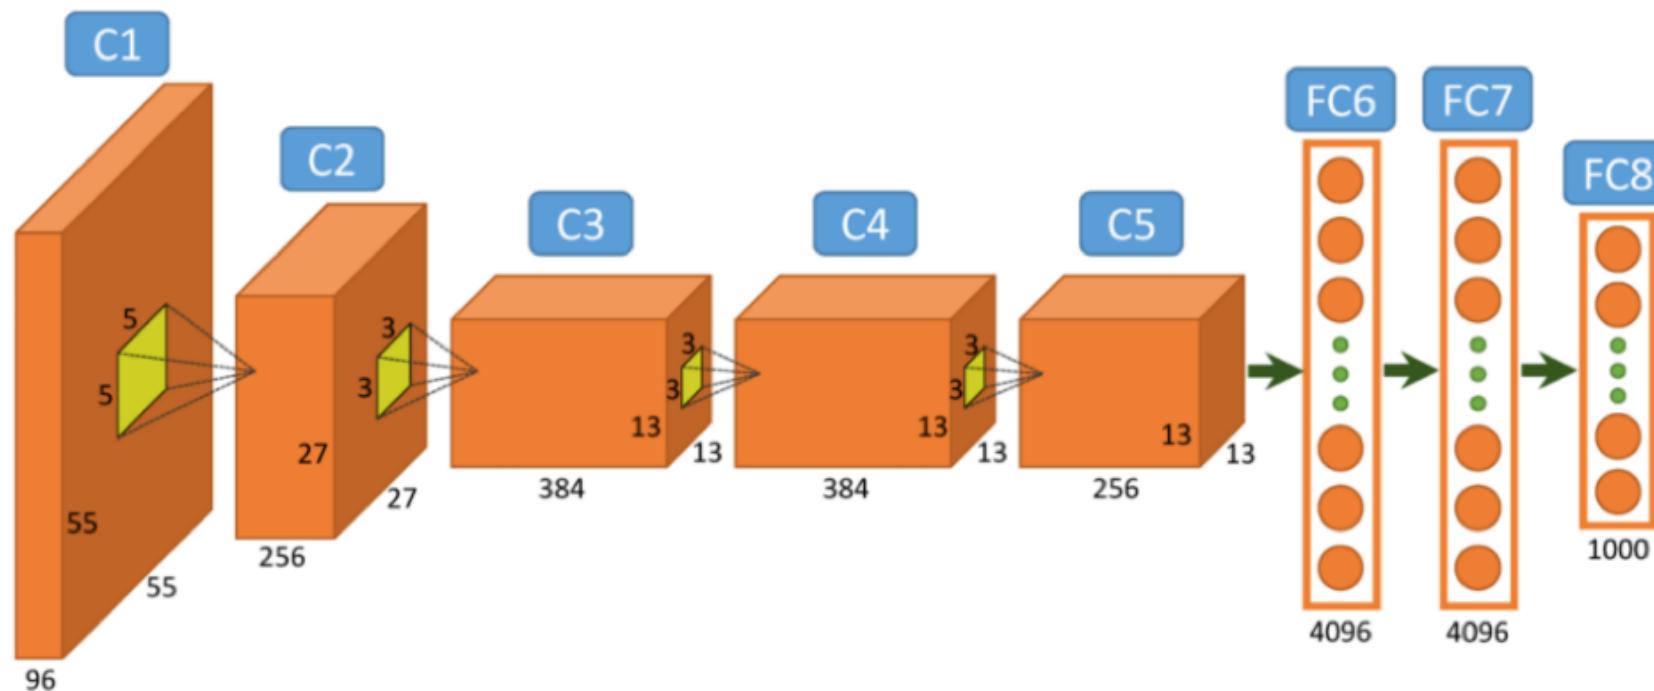
Understanding What Deep Learning Learns

- Distributed representation:
 - e.g., predicting red truck, red car, red bird, green truck, green car, green bird, blue truck, blue car, blue bird
 - Two layers: one layer indicates color (3 nodes) and other layer indicates object (3 nodes)
- Goal: learn “right” representation of the data
- Generally, more convolution layers leads to learning more complex features



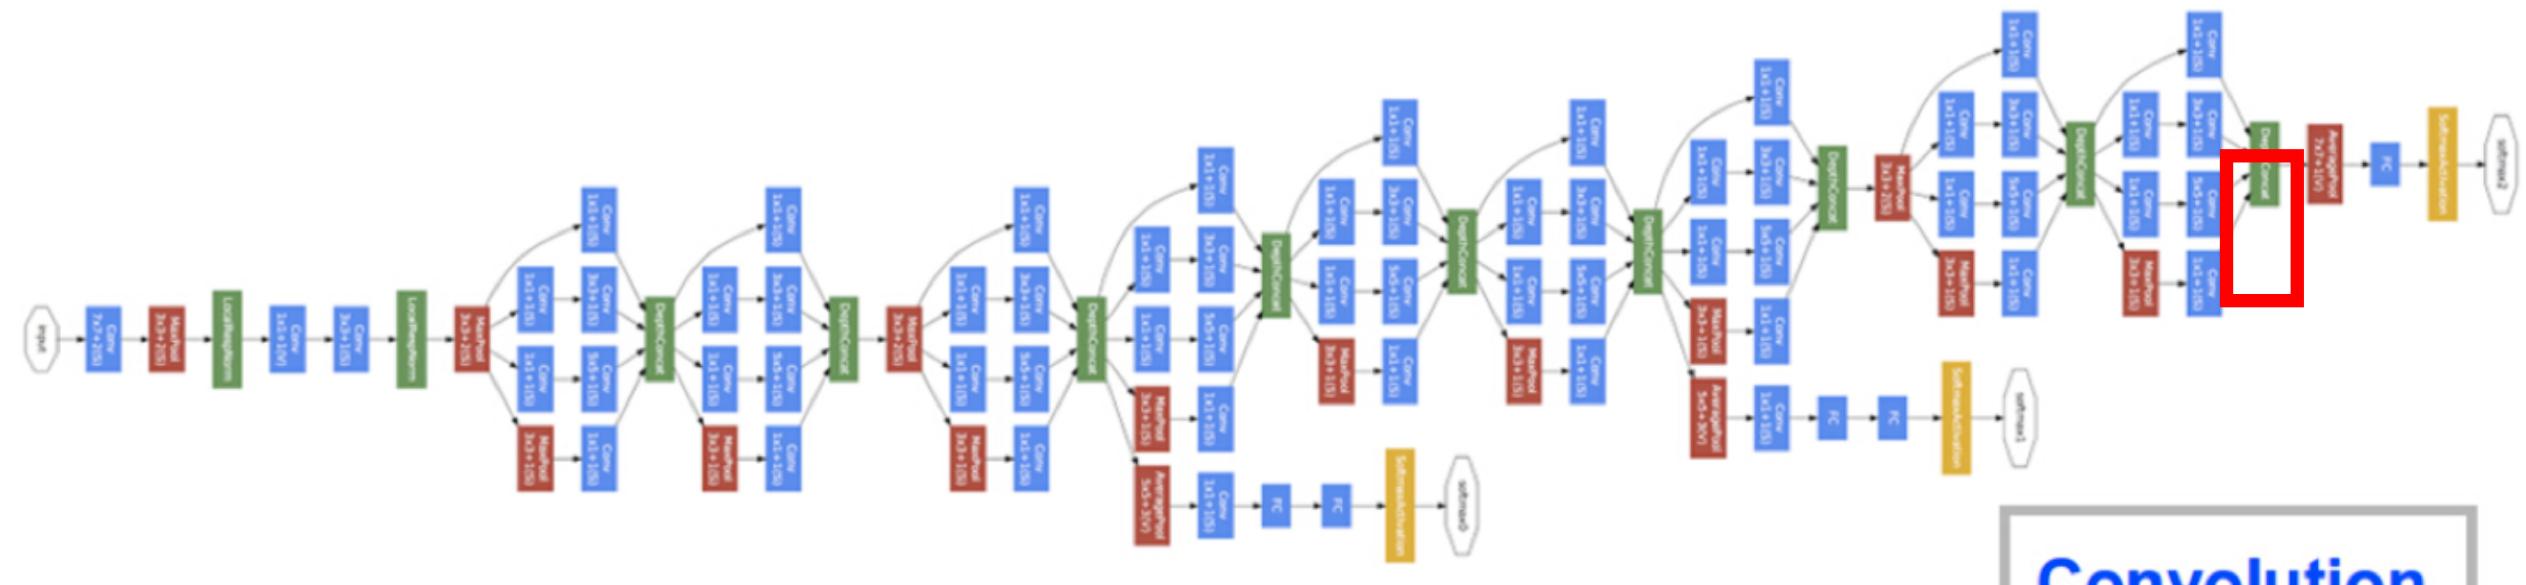
AlexNet Deep Features

- What is the dimensionality of the fc6 feature?
- What is the dimensionality of the fc7 feature?



GoogleNet (Inception) Deep Features

- What is the dimensionality of the inception features?



Today's Topics

- Revisiting History of Modern Neural Networks and “Deep Learning”
- Convolutional Neural Networks (CNN)
- Deep Features
- Lab